

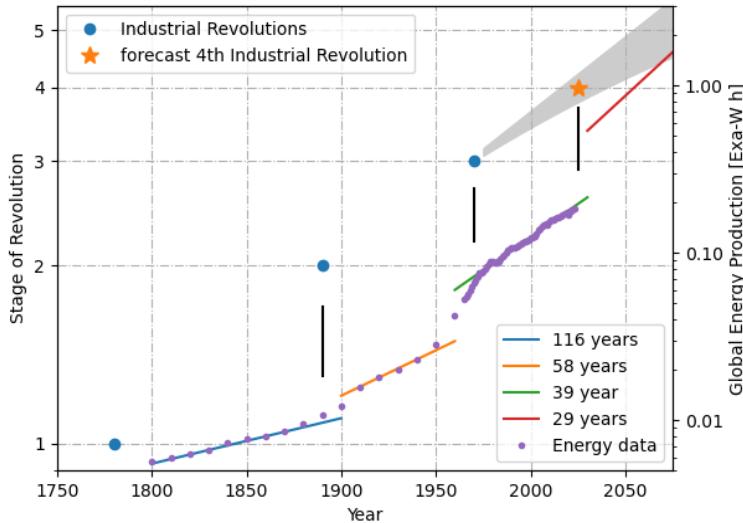
Sensors-to-Edge-to-HPC ecosystem

In flight data analysis for autonomous control systems

Ryan Coffee / Sr. Research Scientist / LCLS+TID+PULSE

October 21, 2025

The Revolution is Now... Ubiquitous Distributed Compute



In the midst of a technology jump

- Greater than 2x jump in global energy consumption, borne by the technically advanced nations
- I bet \$20 on an 8x jump in the US
- Doubling time drops from 39 to **29 years**.



Deployable Adaptivity

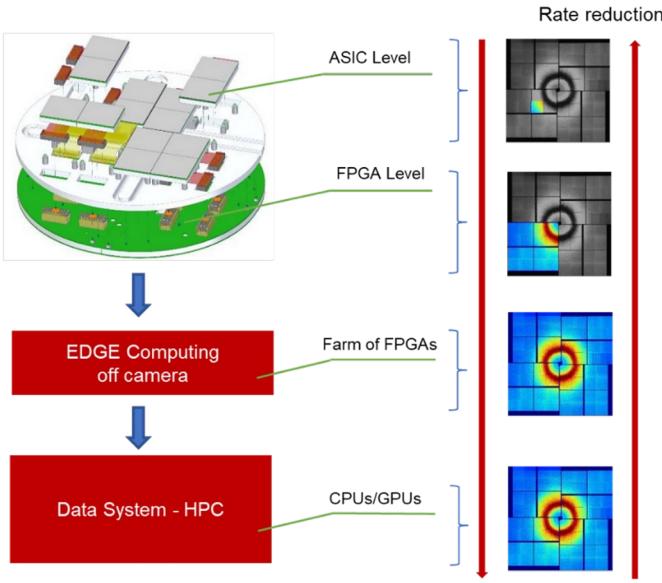
Low-power context-aware inference = **on-device autonomy**

We're not saving energy, just reducing energy per inference and moving to **ubiquitous, scaled, deployed, and continuous training and inference**

The Revolution is Now... Computational Ecosystem

Exemplar Systems of Systems

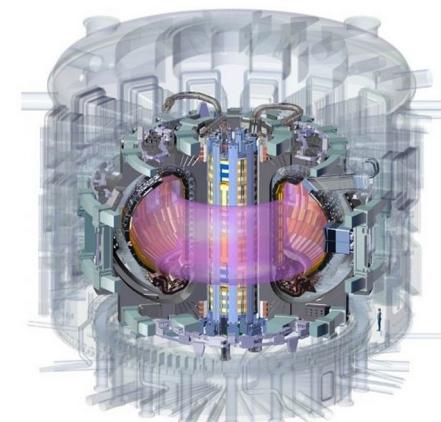
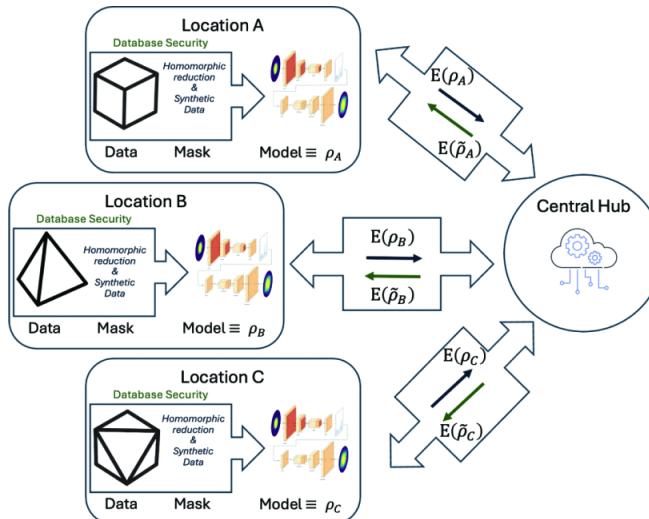
- Interconnected Light Sources
- Interconnected Fission and Fusion Reactors



Courtesy of Jana Thayer
SLAC

Inherently Multi-scale

- $\mu\text{-ms}$ latency decisions
- Minutes-hours scale “run” evolution
- Days-months scale “campaign” forecasting



General Atomics – DIII-D

R. Archibald et al., "Privacy Preserving Federated Learning ...," 2024 IEEE BigData 2024, p. 4132, doi: 10.1109/BigData62323.2024.10825977.

Autonomous Systems of Systems

At the light source

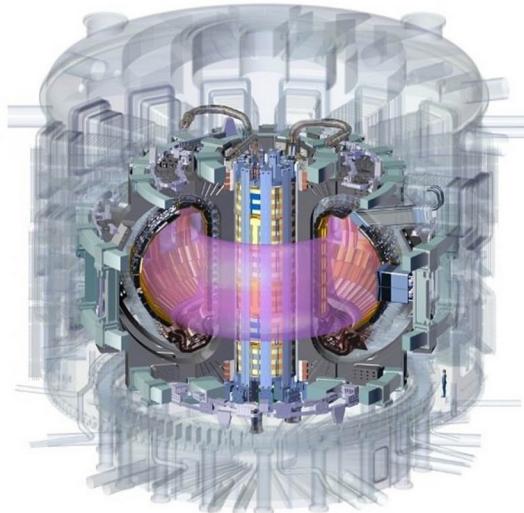
- 1MPix at 1MSps rates
- 1Byte/sample = 1TBps



SLAC Urgency

At the tokamak

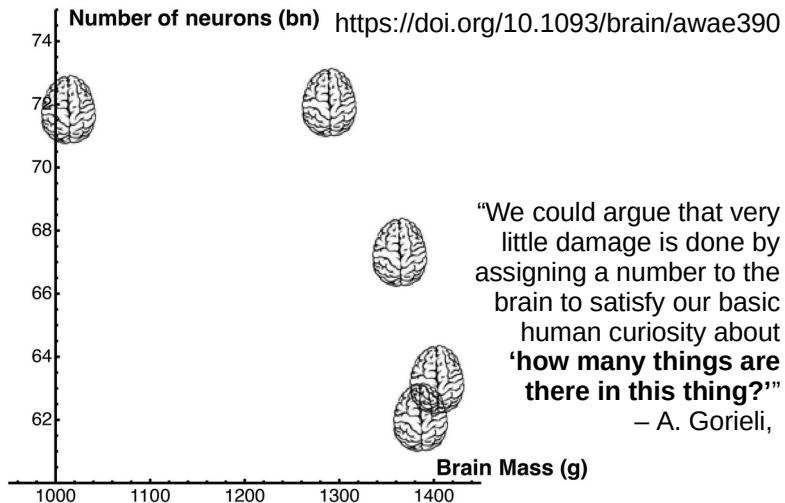
- 1k channels at GSps rates
- 1Byte/sample = 1TBps



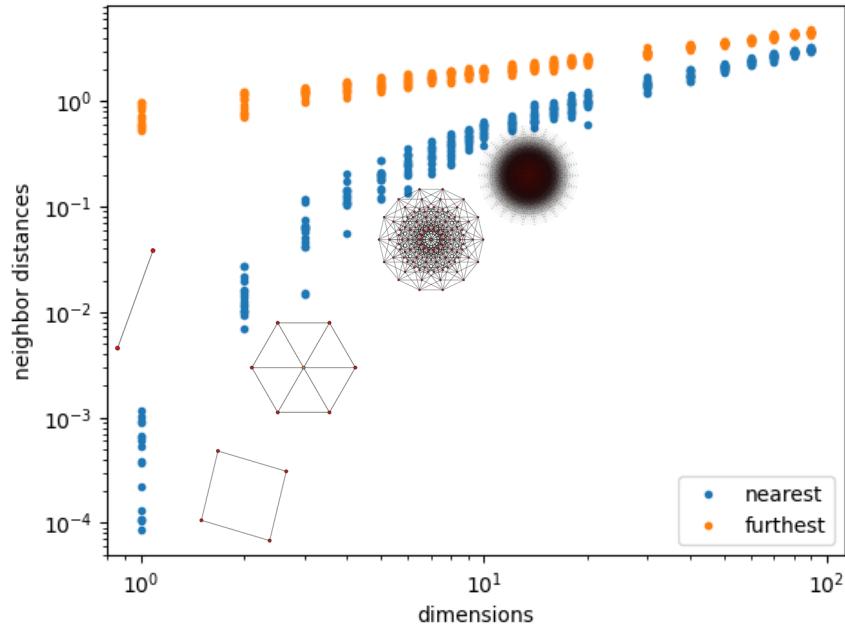
Is AGI actually our goal?

A word for dense information

- Three **orthogonal shadows** are likely better than 13 coupled
- **Interpolation evaporates** in high dimensions
- Reducing the chatter among information channels **reduces extraneous network traffic**
- Correlation dilution across channels increases the cross-section for **information leakage**.



"We could argue that very little damage is done by assigning a number to the brain to satisfy our basic human curiosity about '**how many things are there in this thing?**'"
— A. Goriely,



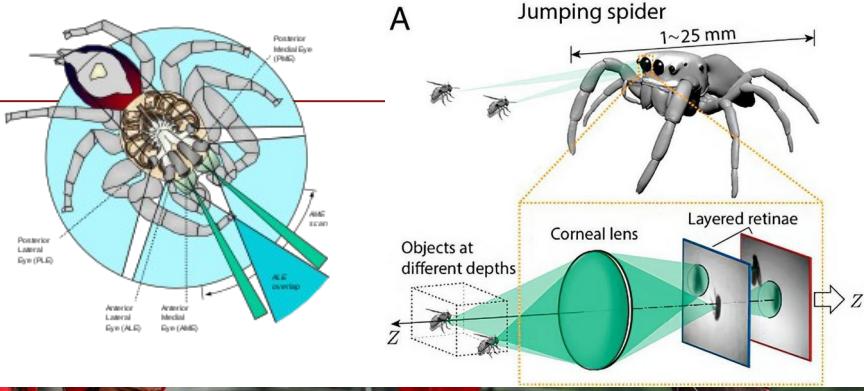
All you need is Edge

Hardware and wetware work in unison

ML in Science is predominantly acceleration of known interpretation

Let's design for Jumping Spider Specificity and efficiency...

... We need an AI Pit Crew

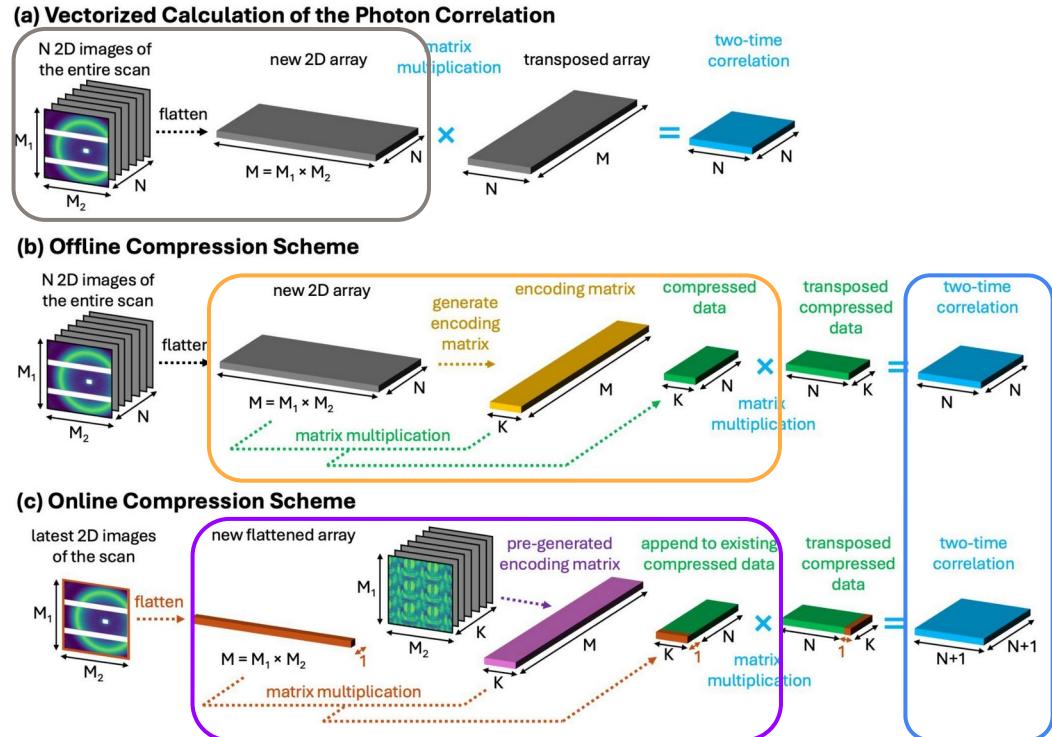


Domain and Detector specific

Generalizable algorithms – what users **need**, not what they **want**

Recast offline for streaming – **offline encoding** explores **information sufficiency**

FPGA/ASIC encoding opens a new can of worms, e.g. **firmware is meta-data**



Stremper *et al.*, "Homomorphic data compression for real time photon correlation analysis," Opt. Express 33, 12059-12070 (2025)

A case study – The Cookie Box

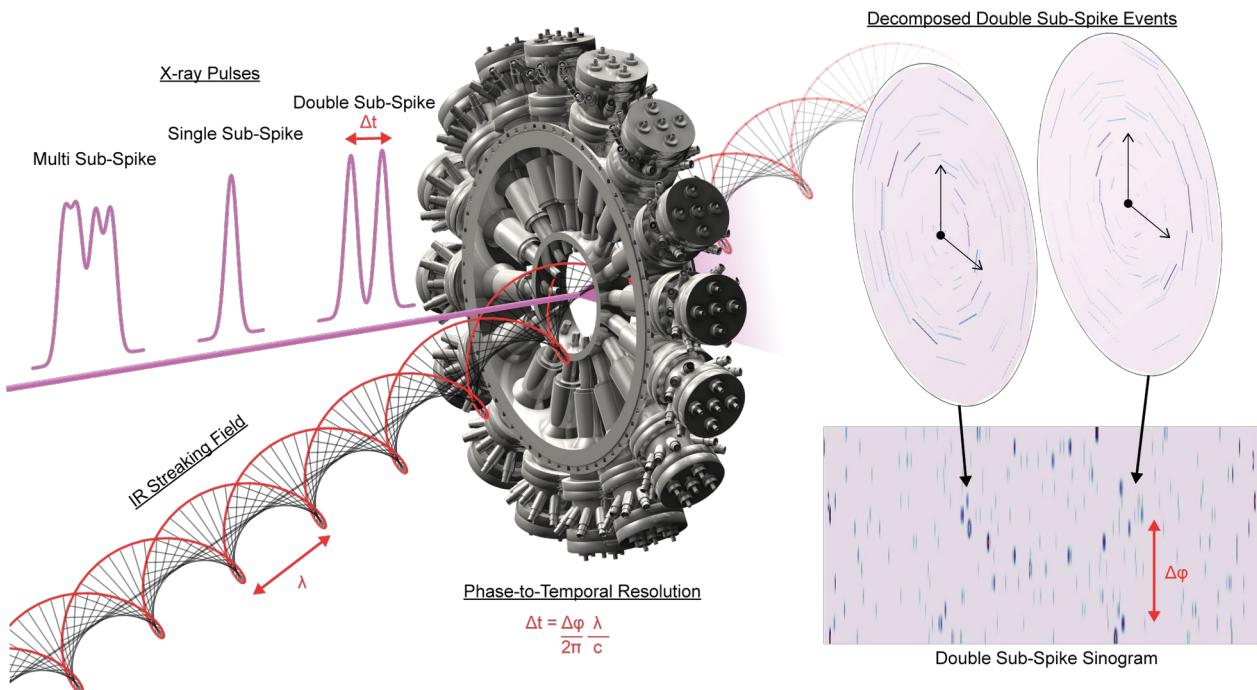
Caught with our hands in the...

X-ray Laser is a **wild ride**, just like the
big wide world

Nature uses **context-dependent**
reasoning/responses

She even pricks ears and shifts gaze
as context evolves

Our detectors should do the same



Autonomous Streaming Decisions

Modular parallel streams

- Denoiser takes longer than 23M parameter Regressor... **Trix** are for FPGAs.
- Motivates **direct coupling of different hardware**
- S3AI working with vendors (Groq and Cornami for now) to **open the communication** to their chips

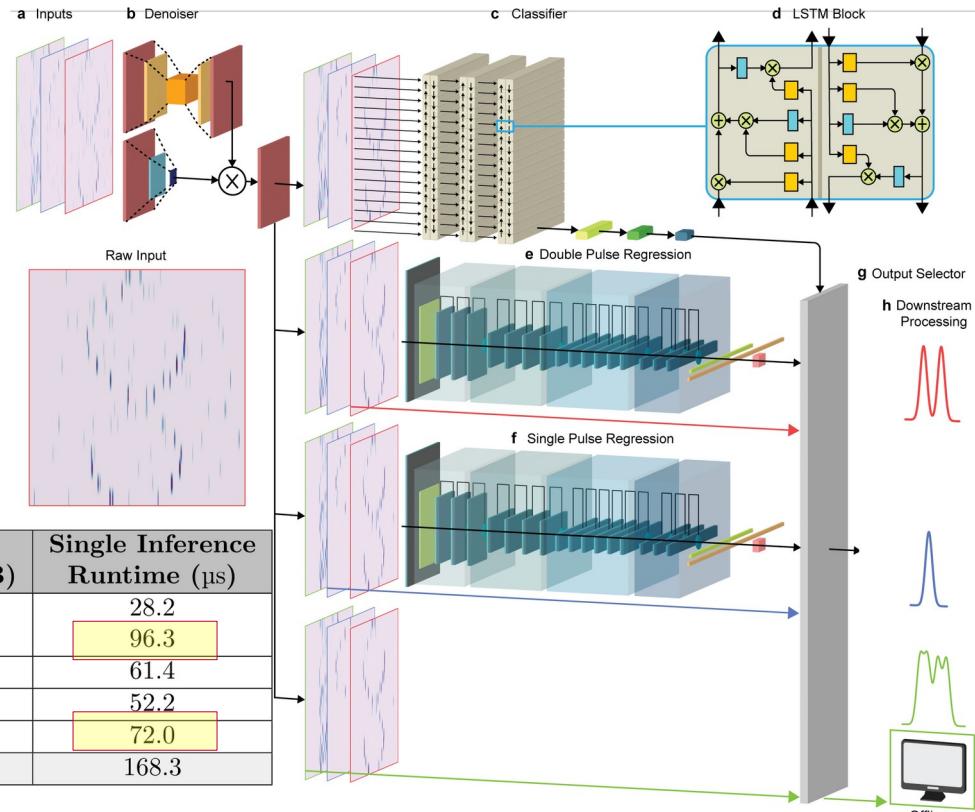


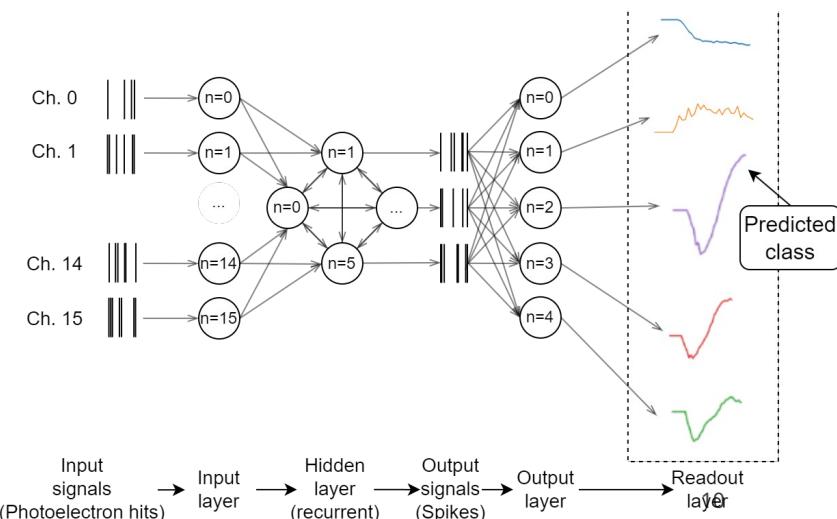
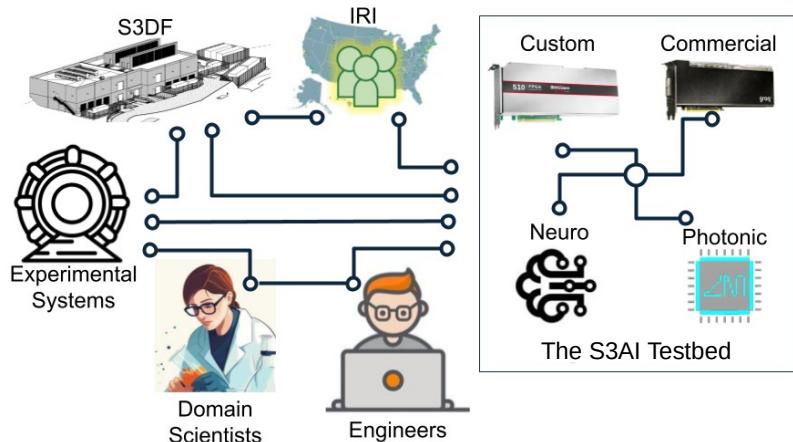
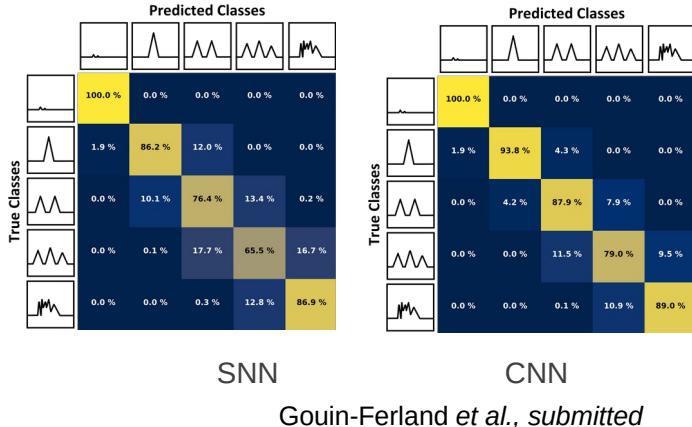
Table 2 DCIPHR parameters and runtime

Model (Identifier)	# Parameters	Parameter Memory (MB)	Single Inference Runtime (μ s)
Denoiser (1)	Zero Classifier (1a)	70,345	28.2
	Autoencoder (1b)	46,529	96.3
Classifier (2)	1,458,597	5.83	61.4
Single Pulse ϕ Regression (3)	12,196,240	48.78	52.2
Double Pulse $\Delta\phi$ Regression (4)	23,330,400	93.32	72.0
Totals	37,102,111	148.40	168.3

A Testbed of our Own

S3AI for heterogeneous exploration

- Simulating direct conversion of time-series into spiking binary signals
- Comparing traditional approaches to **Spiking Neural Networks** ... just as Nature intended!



Taking Control of Our Destiny

Autonomous Control Signals

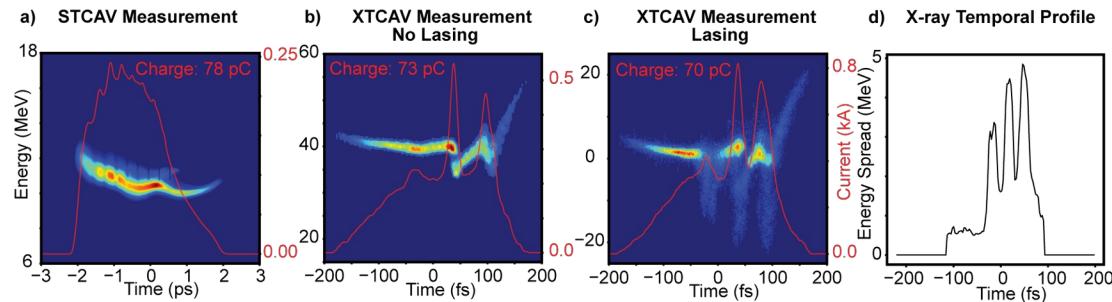
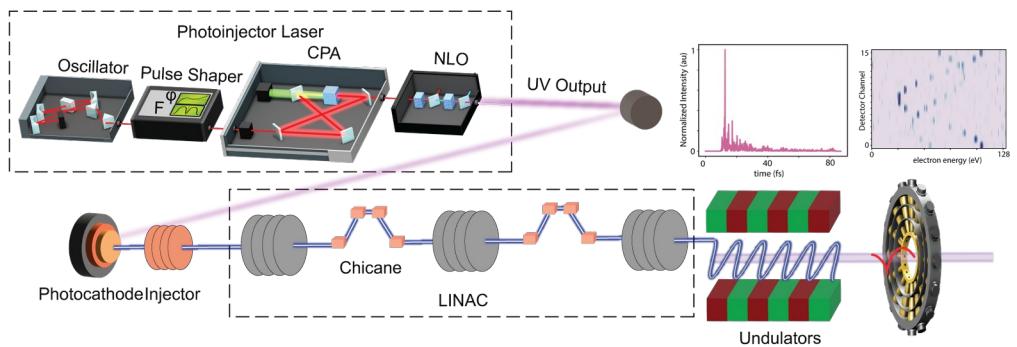
Real-time diagnostics **do not need** to be human-readable

But humans know *how* to read the important minor features

Compression = reduction to those features

Multi-modal signal interpretation

- CookieBox for fast classification
- Xtcav/Passive Streaker for constraint
- Mfps distributed mutual information
- **Control feedback to injector laser shaper**

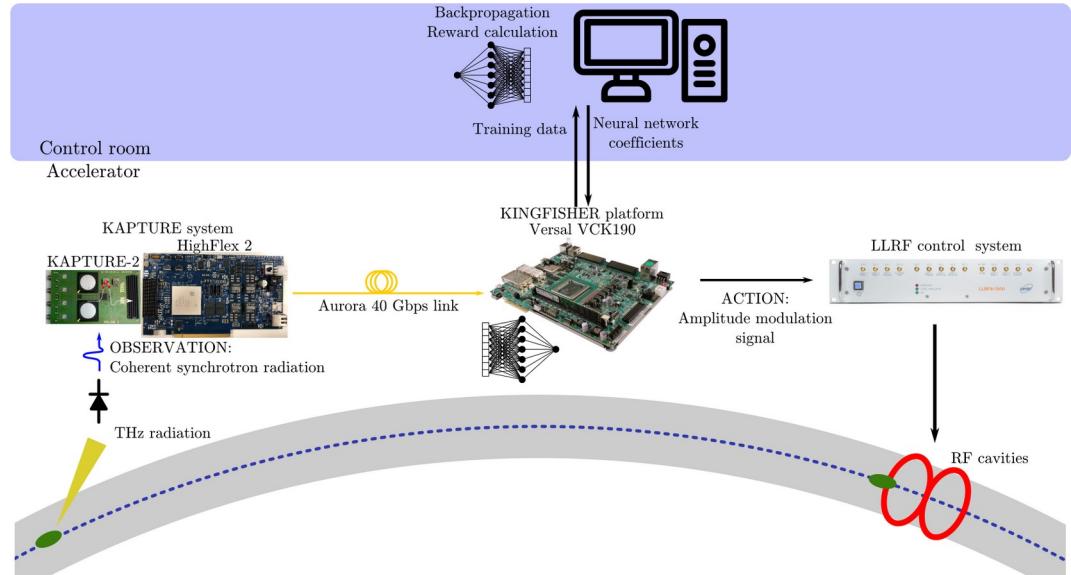


Actual Self-Adaptation – Even on FPGA

Real-time learning +
Ultrafast inference =
Autonomous Revolution

Context control = Agency for AI

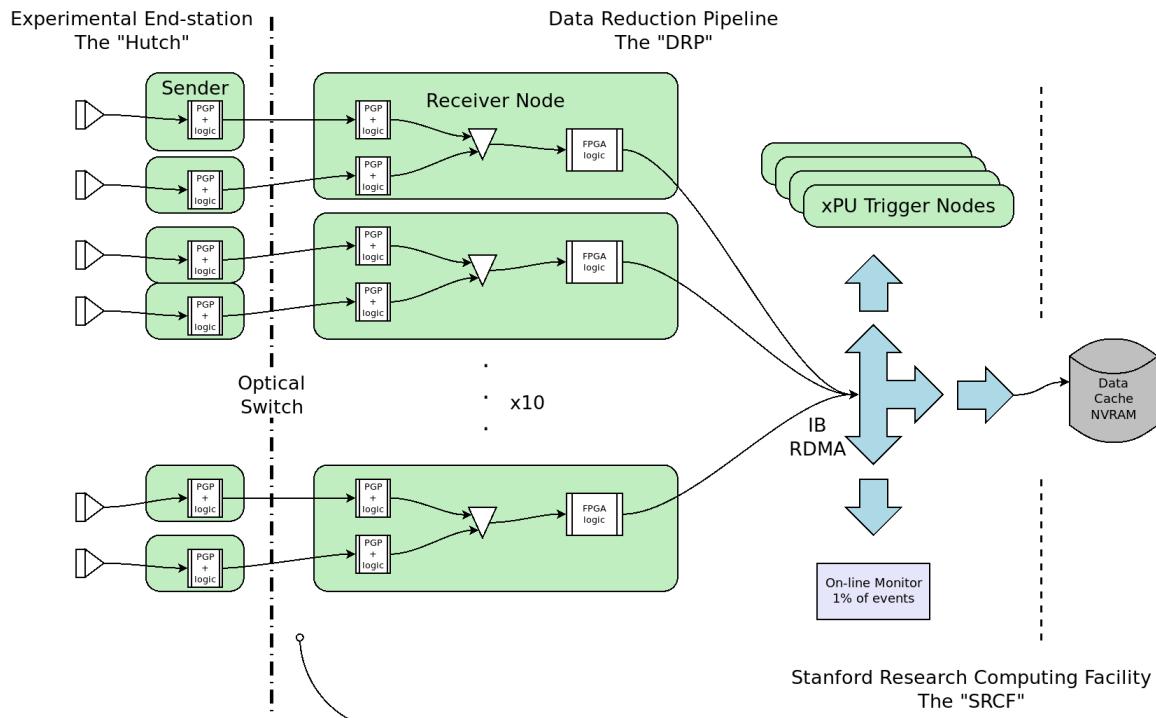
- SNL allows for weight updates in real-time on FPGA
- Weights are trained on remote LCF e.g. Frontier ALCC project LRN045 (DIII-D forecasting)
- But honest autonomy requires bi-directional **real-time exchange** between the sensors, the controls, and the LCF.



L. Scomparin *et al.*, “Preliminary results on the reinforcement learning-based control of the microbunching instability” (2024) 15th IPAC doi:10.18429.JACoW-IPAC2024-TUPS61

Autonomous Decisions – The Light Source Ecosystem

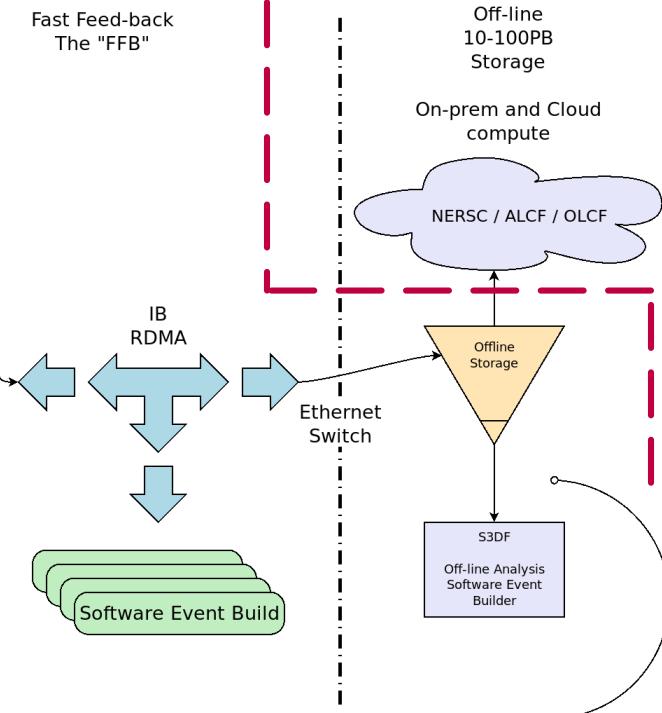
The Data Flow Reality



100s of meters of optical fiber between The Hutch and The DRP.
Event Building is a luxury that users take for granted

LCLStream crosses Blood—Brain Barrier

Fast Feed-back The "FFB"



Autonomous Decisions – The Fusion Ecosystem

At the tokamak (eventually)

- 1k channels at GSps rates
- 2Byte/sample = 2TBps
- Get smart about mutual information

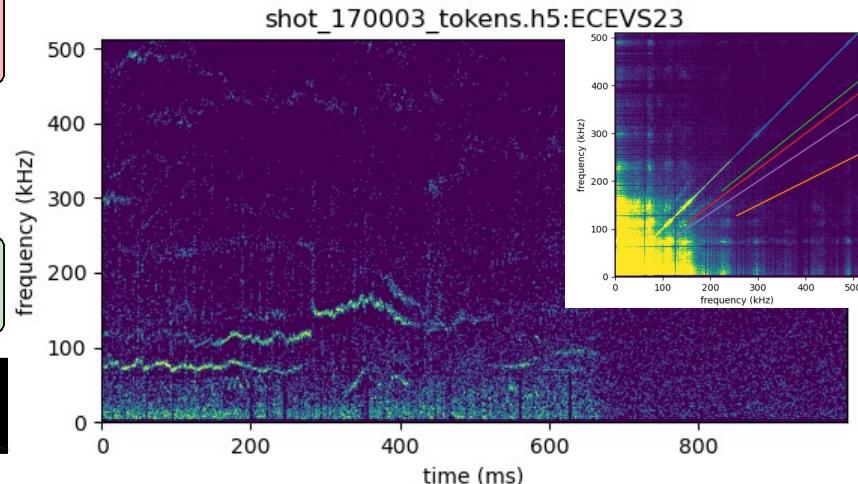
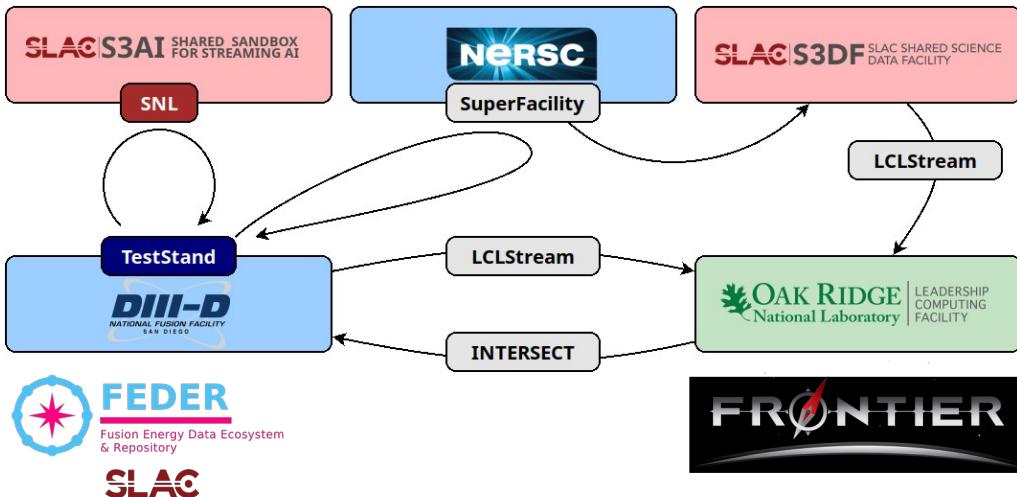


Title:
Principal
Investigator:
Co-investigators:
ALCC Allocation:
Site(s):

Real-Time Adaptive Disruption Forecasting at DIII-D
Ryan Coffee (SLAC National Accelerator Laboratory)

David Rogers (ORNL-NCCS)

Oak Ridge Leadership Computing Facility (OLCF)



Triumvirate for the Win

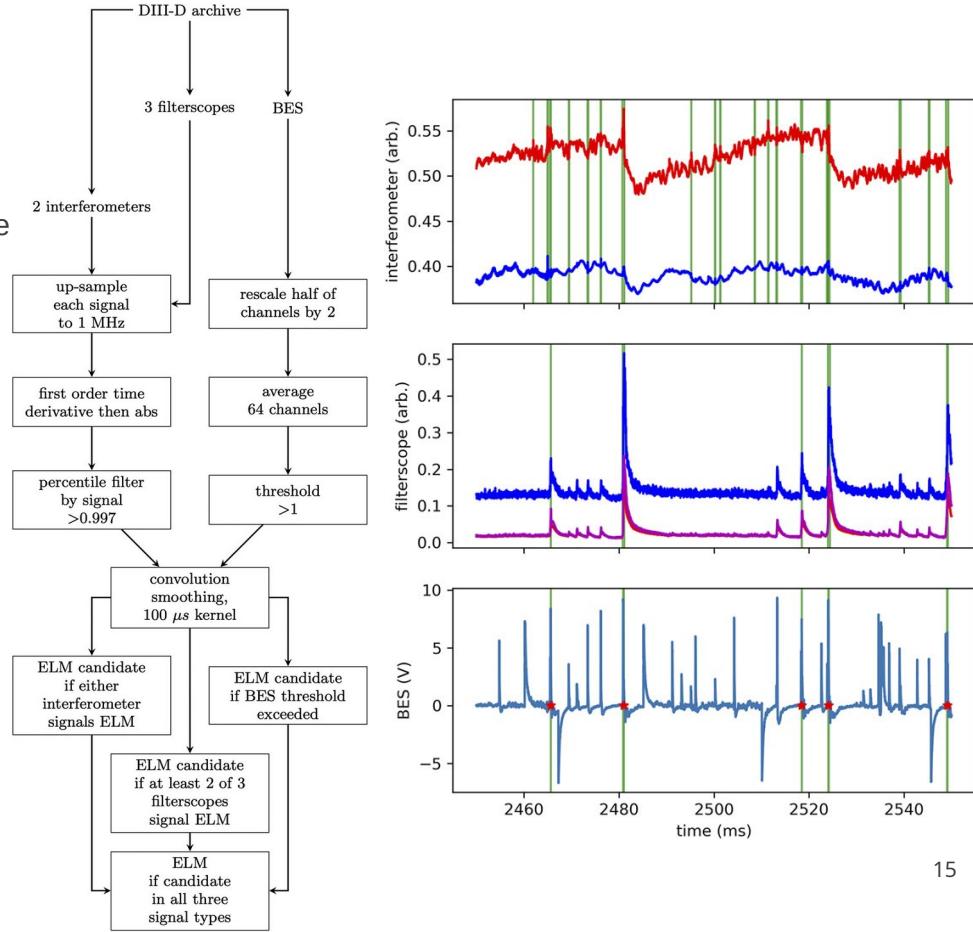
Decision Support Diagnostics

The DIII-D ELM identifier is itself **multi-modal**

- BES is an imaging modality highest sample rate
- Filterscopes are point detection but uniquely see ELMs
- Interferometers are ultrasensitive belt & suspenders

Incarnation of a mixture of **different** experts

Finn H. O'Shea, et al., "Automatic identification of edge localized modes in the DIII-D tokamak" APL Mach. Learn. (2023) 1, (2) 026102 <https://doi.org/10.1063/5.0134001>

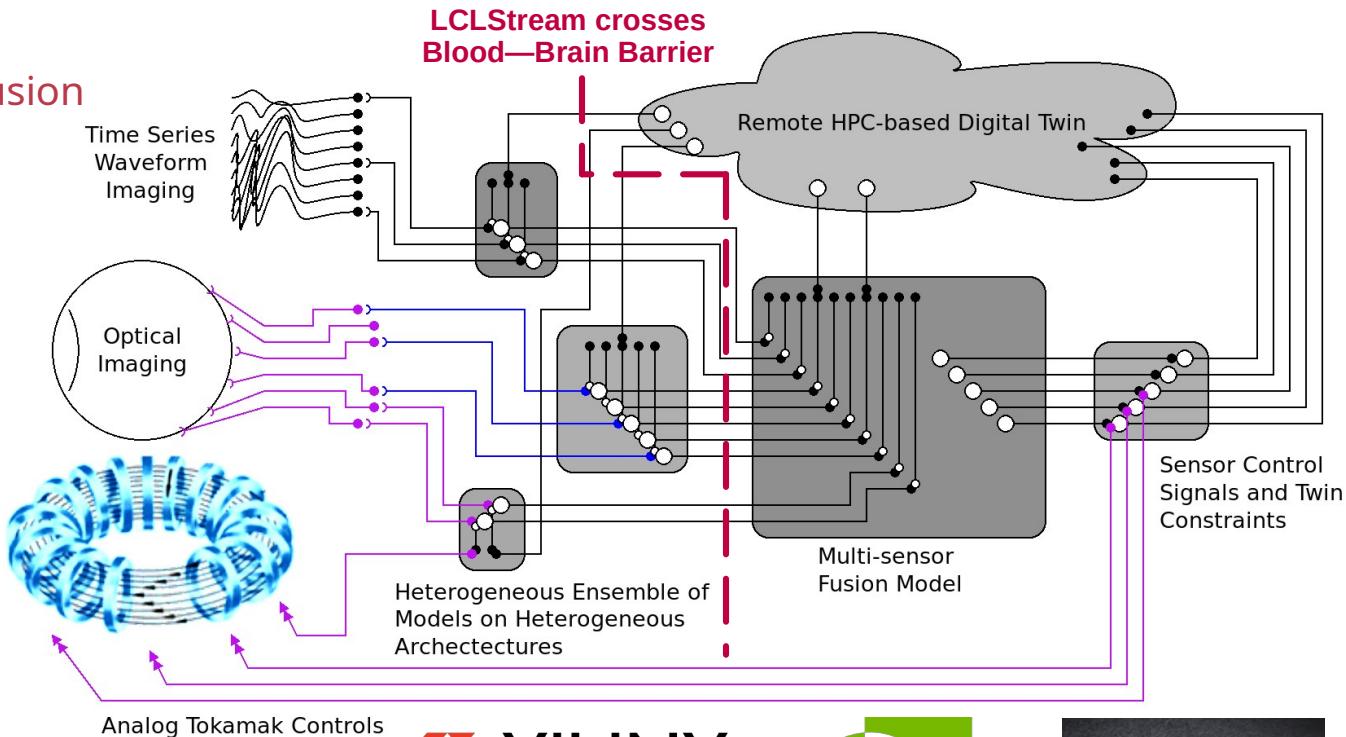


An Important Destiny to Control

Orthogonal Views on Fusion

Breaking through silos, linking across Labs, and partnering with industry

- LCLStream model for **Edge-to-Exa streaming**
- INTERSECT enables geographically **remote autonomous control**
- Embrace industry collaboration for **direct coupling of sensing to inference acceleration**



Fitting the Pieces Together for the Ecosystem

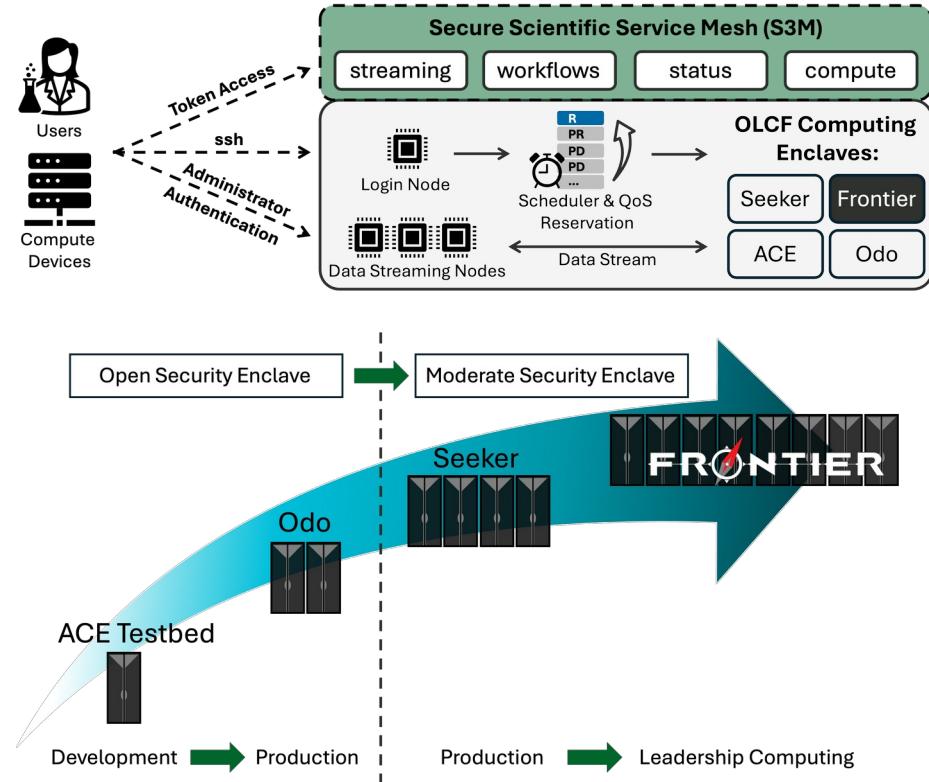
Shovel-ready Components

- Focus on what we **have** rather than what we **want**
- **Interconnect modules** rather than boil the ocean
- Must **break down rivalries**
- This is the time to **burrow under silo walls**



Title:	Real-Time Adaptive Disruption Forecasting at DIII-D
Principal Investigator:	Ryan Coffee (SLAC National Accelerator Laboratory)
Co-investigators:	David Rogers (ORNL-NCCS)
ALCC Allocation:	Oak Ridge Leadership Computing Facility (OLCF)
Site(s):	Oak Ridge Leadership Computing Facility (OLCF)
Allocation(s):	350,000 on Frontier
Research Summary:	<p>This project will extend an existing statically trained machine-learning based disruption prediction model for tokamak fusion reactors by leveraging a meta-learning method for fast optimization of the plasma state forecasting model and also leveraging an encoder/decoder model that accommodates a dynamic quantization scheme. The quantization will be optimized to explore the space of model and feature encoding. This exploration will inform decisions about reduced fidelity diagnostics that would remain sufficient for a reactor regime tokamak state prediction consistent for energy on the grid.</p>

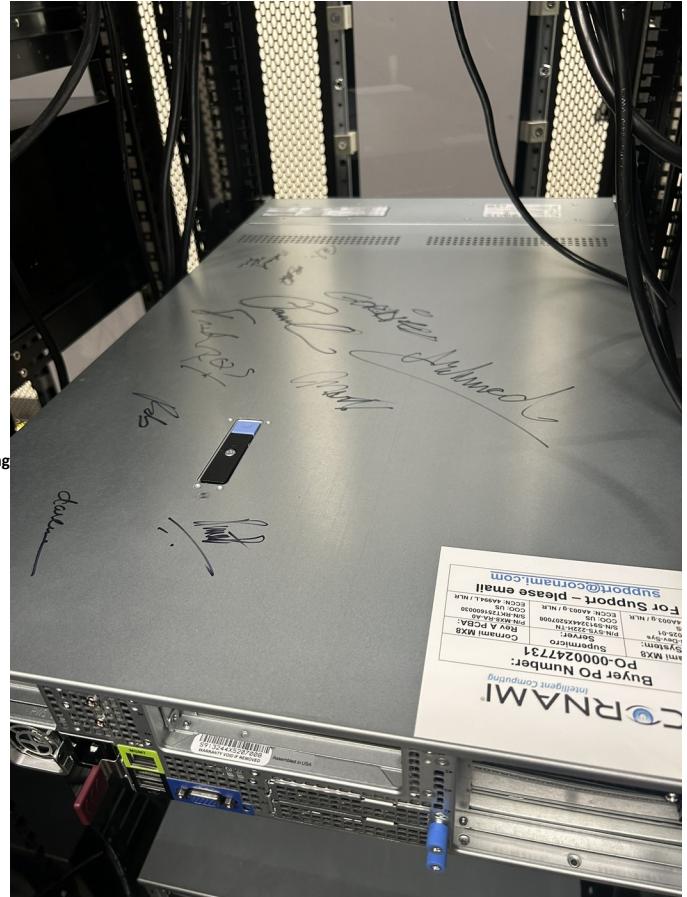
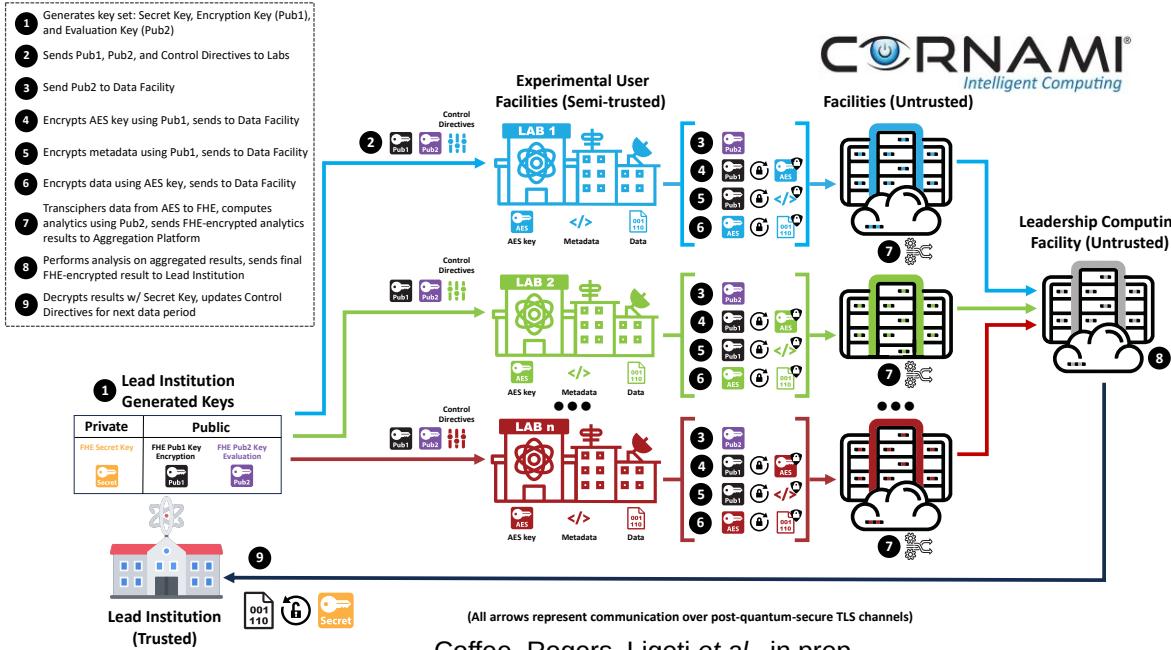
2025 ASCR Leadership Computing Challenge Award



Future-proofing the Ecosystem

Inherently Multi-Party

- Security with **Shared Resources/Orthogonal views**
 - Interconnecting **secure enclaves** with open cloud
 - Design for “**Collaborating Competitors**”



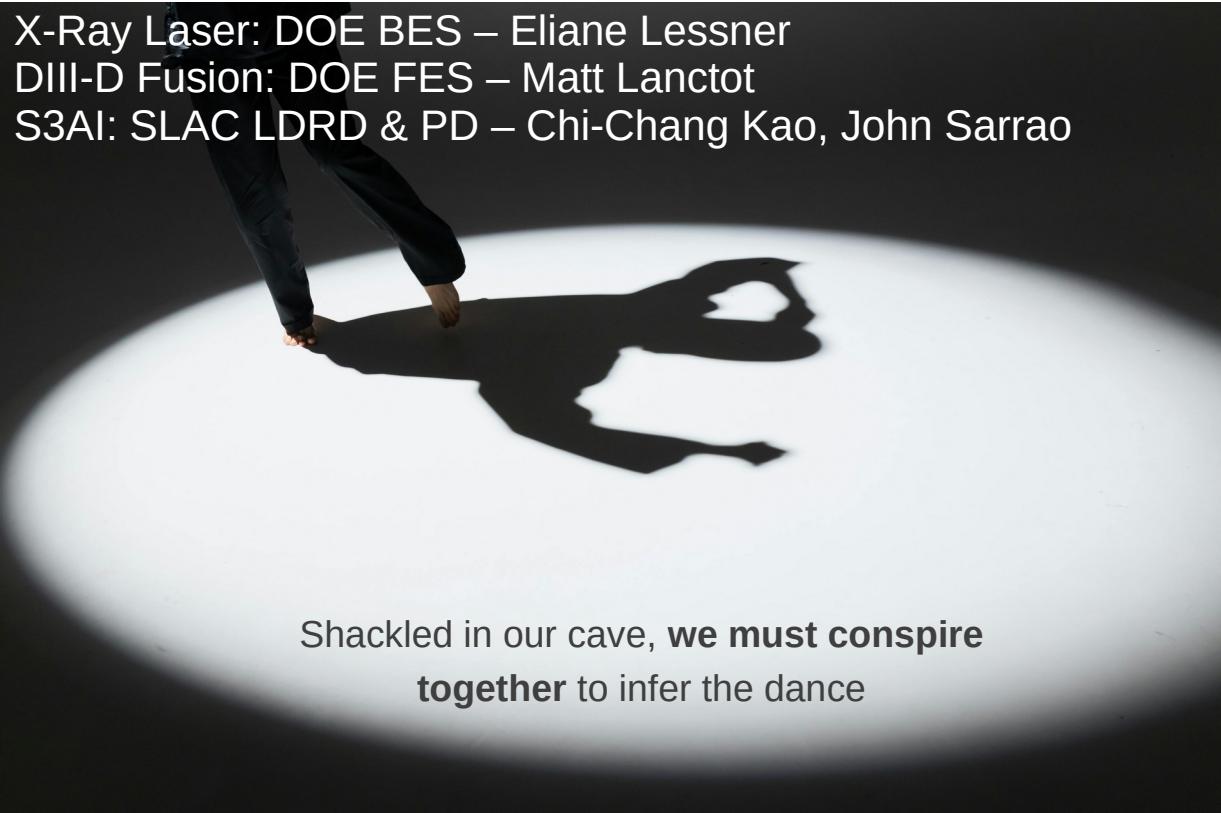
How I learned to stop worrying...

Dance of Shadows

Although beautiful, a single shadow does not well represent the dancer

Multi-modal signal interpretation

- Every sensor is responsible for its own **brutal parsimony**
- Nature abhors redundant arrays
- Autonomy requires a rich, diverse, and efficient sensor environments
- It's all for naught if you have to deliberate every move ...
latency really matters



Shackled in our cave, **we must conspire together** to infer the dance

The background of the slide features a stage set against a dark background. The stage is illuminated by several bright spotlights from above, creating sharp beams of light. There are four large rectangular structures resembling windows or panels on the left and right sides. In the center, there's a stack of various materials, including what looks like stone blocks and wooden crates, with some boxes labeled "MATERIALS".

Thank You!

Now let's get crackin'

coffee@slac.stanford.edu

