



# **Neural Operators: A Scalable Framework for AI-Driven Scientific Discovery**

Jean Kossaifi


LSSDA Workshop

SLAC National Accelerator Laboratory

October 2025

# Moving beyond computer vision and language

Capturing the continuum



# Realistic Climate Simulation is a Computational Grand Challenge

Capturing Fine-Scales is Too Expensive (currently)





Figure adapted from: Schneider, T., Teixeira, J., Bretherton, C. et al. "Climate goals and computing the future of clouds". *Nature Climate Change* 7, 3–5 (2017)

# Speedup leads to better probabilistic estimates

Large ensembles in seconds





# **Methods for Scientific Computing and Engineering**

# Capturing the continuum: neural operators

Train and Inference at any discretization




Neural Network

Input and output at fixed resolution



Neural Operator







Input and output at any points in domain



# Capturing the continuum: neural operators

- Query at any point in the domain
- Decouple resolution and number of parameters
- Converges upon mesh refinement to a limit (continuum).


Mesh refinement



Converging solution

# A General Framework for Neural Operators

Integral Linear Operator as a Core Building Block



# Fourier Neural Operator

Fourier transform for global convolution

Integral linear operator

$$\int \kappa(x, y) v(y) dy$$


Convolution operator  
(special case of integral operator)

$$\int \kappa(x - y) v(y) dy$$


Perform convolution in Fourier domain

$$\mathcal{F}^{-1}( \mathcal{F}(\kappa) . \mathcal{F}(v) )$$

Learn weights  $R$  in Fourier Domain




$$R := \mathcal{F}(\kappa)$$



# FNO: Fourier Neural operator


Discretization invariant learning on regular grids

- (Fast) Fourier Transform implements global convolution
- Compose global convolution with non-linearity




# Learning across resolutions with Neural Operators

training data  
(64x64)





prediction  
(256x256)




# **Learning on a Sphere**

# Fourier Transform on a Sphere

Sine/cosine  $\rightarrow$  circular functions (spherical harmonics)






2018-05-01


# AI FOR EXTREME WEATHER FORECASTING

SFNO enables larger ensembles and **better risk assessment**

50 ensemble members



1000 ensemble members



# **Learning on Arbitrary Geometries**

# Graph Neural Operator

$$\int \kappa(x, y) v(y) dy$$


↓

$$v_l(x) = \int_{B_r(x)} \kappa(x, y) v_{l-1}(y) dy$$

↓

Layer l

Integrate over local ball,  
centered on point x,  
with radius r,



# Scaling Neural Operators

# Scale is Often the Best Path to Better Skill

However, experimentation at scale is difficult




Image from: Scaling Laws for Neural Language Models, Kaplan et al, 2020

# Scaling up


Experimentation at scale is difficult

- Need specialized implementation (e.g. Spherical harmonics)
- Need to handle heterogeneous data/geometries:
  - Batching is hard
  - Computationally expensive
- I/O and storage become an issue as resolution increases
- Need for various types of parallelism (Data Parallel, Fully-Sharded Data Parallel, Domain Parallelism)

# Scaling up

Experimentation at scale is difficult


- Custom kernels
- Simultaneous data & model parallelism
- Incremental training
- Tensor Factorization



# **Neural Operators in Python**



# Learning Neural Operators in Python



<https://github.com/neuraloperator/neuraloperator>

```
model = FNO(n_modes=(16, 16),  
            in_channels=1,  
            out_channels=1,  
            hidden_channels=32,  
            projection_channel_ratio=2)  
model = model.to(device)
```

```
l2loss = LpLoss(d=2, p=2)  
h1loss = H1Loss(d=2)
```


# Open Science

Model code, checkpoints, massively parallel training code and inference/scoring code

README License

## torch-harmonics

[Overview](#) | [Installation](#) | [More information](#) | [Getting started](#) | [Contributors](#) | [Cite us](#) | [References](#)

 passing |  v0.8.0


### Overview

torch-harmonics implements differentiable signal processing on the sphere. This includes differentiable implementations of the spherical harmonic transforms, vector spherical harmonic transforms and discrete-continuous convolutions on the sphere. The package was originally implemented to enable Spherical Fourier Neural Operators (SFNO) [1].

The SHT algorithm uses quadrature rules to compute the projection onto the associated Legendre polynomials and FFTs for the projection onto the harmonic basis. This algorithm tends to outperform others with better asymptotic scaling for most practical purposes [2].

torch-harmonics uses PyTorch primitives to implement these operations, making it fully differentiable. Moreover, the quadrature can be distributed onto multiple ranks making it spatially distributed.

torch-harmonics has been used to implement a variety of differentiable PDE solvers which generated the animations below. Moreover, it has enabled the development of Spherical Fourier Neural Operators [1].



### Installation


A simple installation can be directly done from PyPi:

```
pip install torch-harmonics
```

README License


## Makani: Massively parallel training of machine-learning based weather and climate models

[Overview](#) | [Getting started](#) | [More information](#) | [Contributing](#) | [Further reading](#) | [References](#)

 passing

Makani (the Hawaiian word for wind 🌬️) is a library designed to enable the research and development of the next generation of machine-learning (ML) based weather and climate models in PyTorch. Makani was used to train [FourCastNet3](#) [1], [Spherical Fourier Neural Operators \(SFNO\)](#) [2] for weather (FourCastNet2), [Huge ensemble of SFNO \(HENS-SFNO\)](#) [3,4], and [FourCastNet1](#) [5].

Makani is aimed at researchers working on ML based weather prediction. Stable features are frequently ported to the [earth2studio](#) and the [NVIDIA PhysicsNeMo](#) framework. For commercial and production purposes, we recommend checking out these packages.



### Overview

Makani is a research code developed by engineers and researchers at NVIDIA and NERSC for massively parallel training of weather and climate prediction models on 100+ GPUs and to enable the development of the next

# Open Source Software

Key technologies regularly ported to mature frameworks

README Apache-2.0 license Security

## NVIDIA PhysicsNeMo

NVIDIA Modulus has been renamed to NVIDIA PhysicsNeMo

repo status Active license Apache-2.0 code style black

[Nvidia PhysicsNeMo](#) | [Documentation](#) | [Install guide](#) | [Getting Started](#) | [Contributing Guidelines](#) | [License](#)

### What is PhysicsNeMo?

NVIDIA PhysicsNeMo is an open-source deep-learning framework for building, training, fine-tuning and inferring Physics AI models using state-of-the-art SciML methods for AI4science and engineering.

PhysicsNeMo provides python modules to compose scalable and optimized training and inference pipelines to explore, develop, validate and deploy AI models that combine physics knowledge with data, enabling real-time predictions.

Whether you are exploring the use of Neural operators, GNNs, or transformers or are interested in Physics-informed Neural Networks or a hybrid approach in between, PhysicsNeMo provides you with an optimized stack that will enable you to train your models at scale.

Navier Stokes      Maxwell's equations  
 $\rho \frac{D\mathbf{u}}{Dt} = -\nabla p + \rho\mathbf{g} + \mu\nabla^2\mathbf{u}$        $\nabla \times \mathbf{E} = -\frac{\partial \mathbf{B}}{\partial t}$   
Linear Elasticity      Dataset      Wave Equation  
 $\varepsilon = \frac{1}{2}[\nabla\mathbf{u} + (\nabla\mathbf{u})^T]$        $\frac{\partial^2 u}{\partial t^2} = c^2\nabla^2 u$   
Generalizable SciML model

README Apache-2.0 license

## NVIDIA Earth2Studio

Python 3.10 | 3.11 | 3.12 | 3.13 License Apache 2.0 coverage 92% mypy Checked Code Style Black Ruff uv

Earth2Studio is a Python-based package designed to get users up and running with AI Earth system models fast. Our mission is to enable everyone to build, research and explore AI driven weather and climate science.

- Earth2Studio Documentation -

[Install](#) | [User-Guide](#) | [Examples](#) | [API](#)

 Pre-Trained Model Zoo       Cloud Datasources       GPU Accelerated Utilities       Composable APIs

### Quick start

Install Earth2Studio:

```
pip install earth2studio[dlwp]
```

Run a deterministic AI weather prediction in just a few lines of code:

```
from earth2studio.models.px import DLWP
from earth2studio.data import GFS
from earth2studio.io import NetCDF4Backend
from earth2studio.run import deterministic as run

model = DLWP.load_model(DLWP.load_default_package())
ds = GFS()
io = NetCDF4Backend("output.nc")
```

# Thank you!

Nikola Kovachki

Daniel Leibovici

Anima Anandkumar

Jan Kautz

David Pitt

Kamyar Azzizadensheli

Boris Bonev

Thorsten Knuth

Mike Pritchard

Robert Joseph George

Valentin Duruisseaux

Zongyi Li

Jiawei

Md Ashiqur Rahman

Miguel Liu-Schiaffini

Julius Berner

Colin White

Chris Choy

Jaideep Pathak

Yannis Panagakis

... and many others!



Thank you!

X @JeanKossaifi