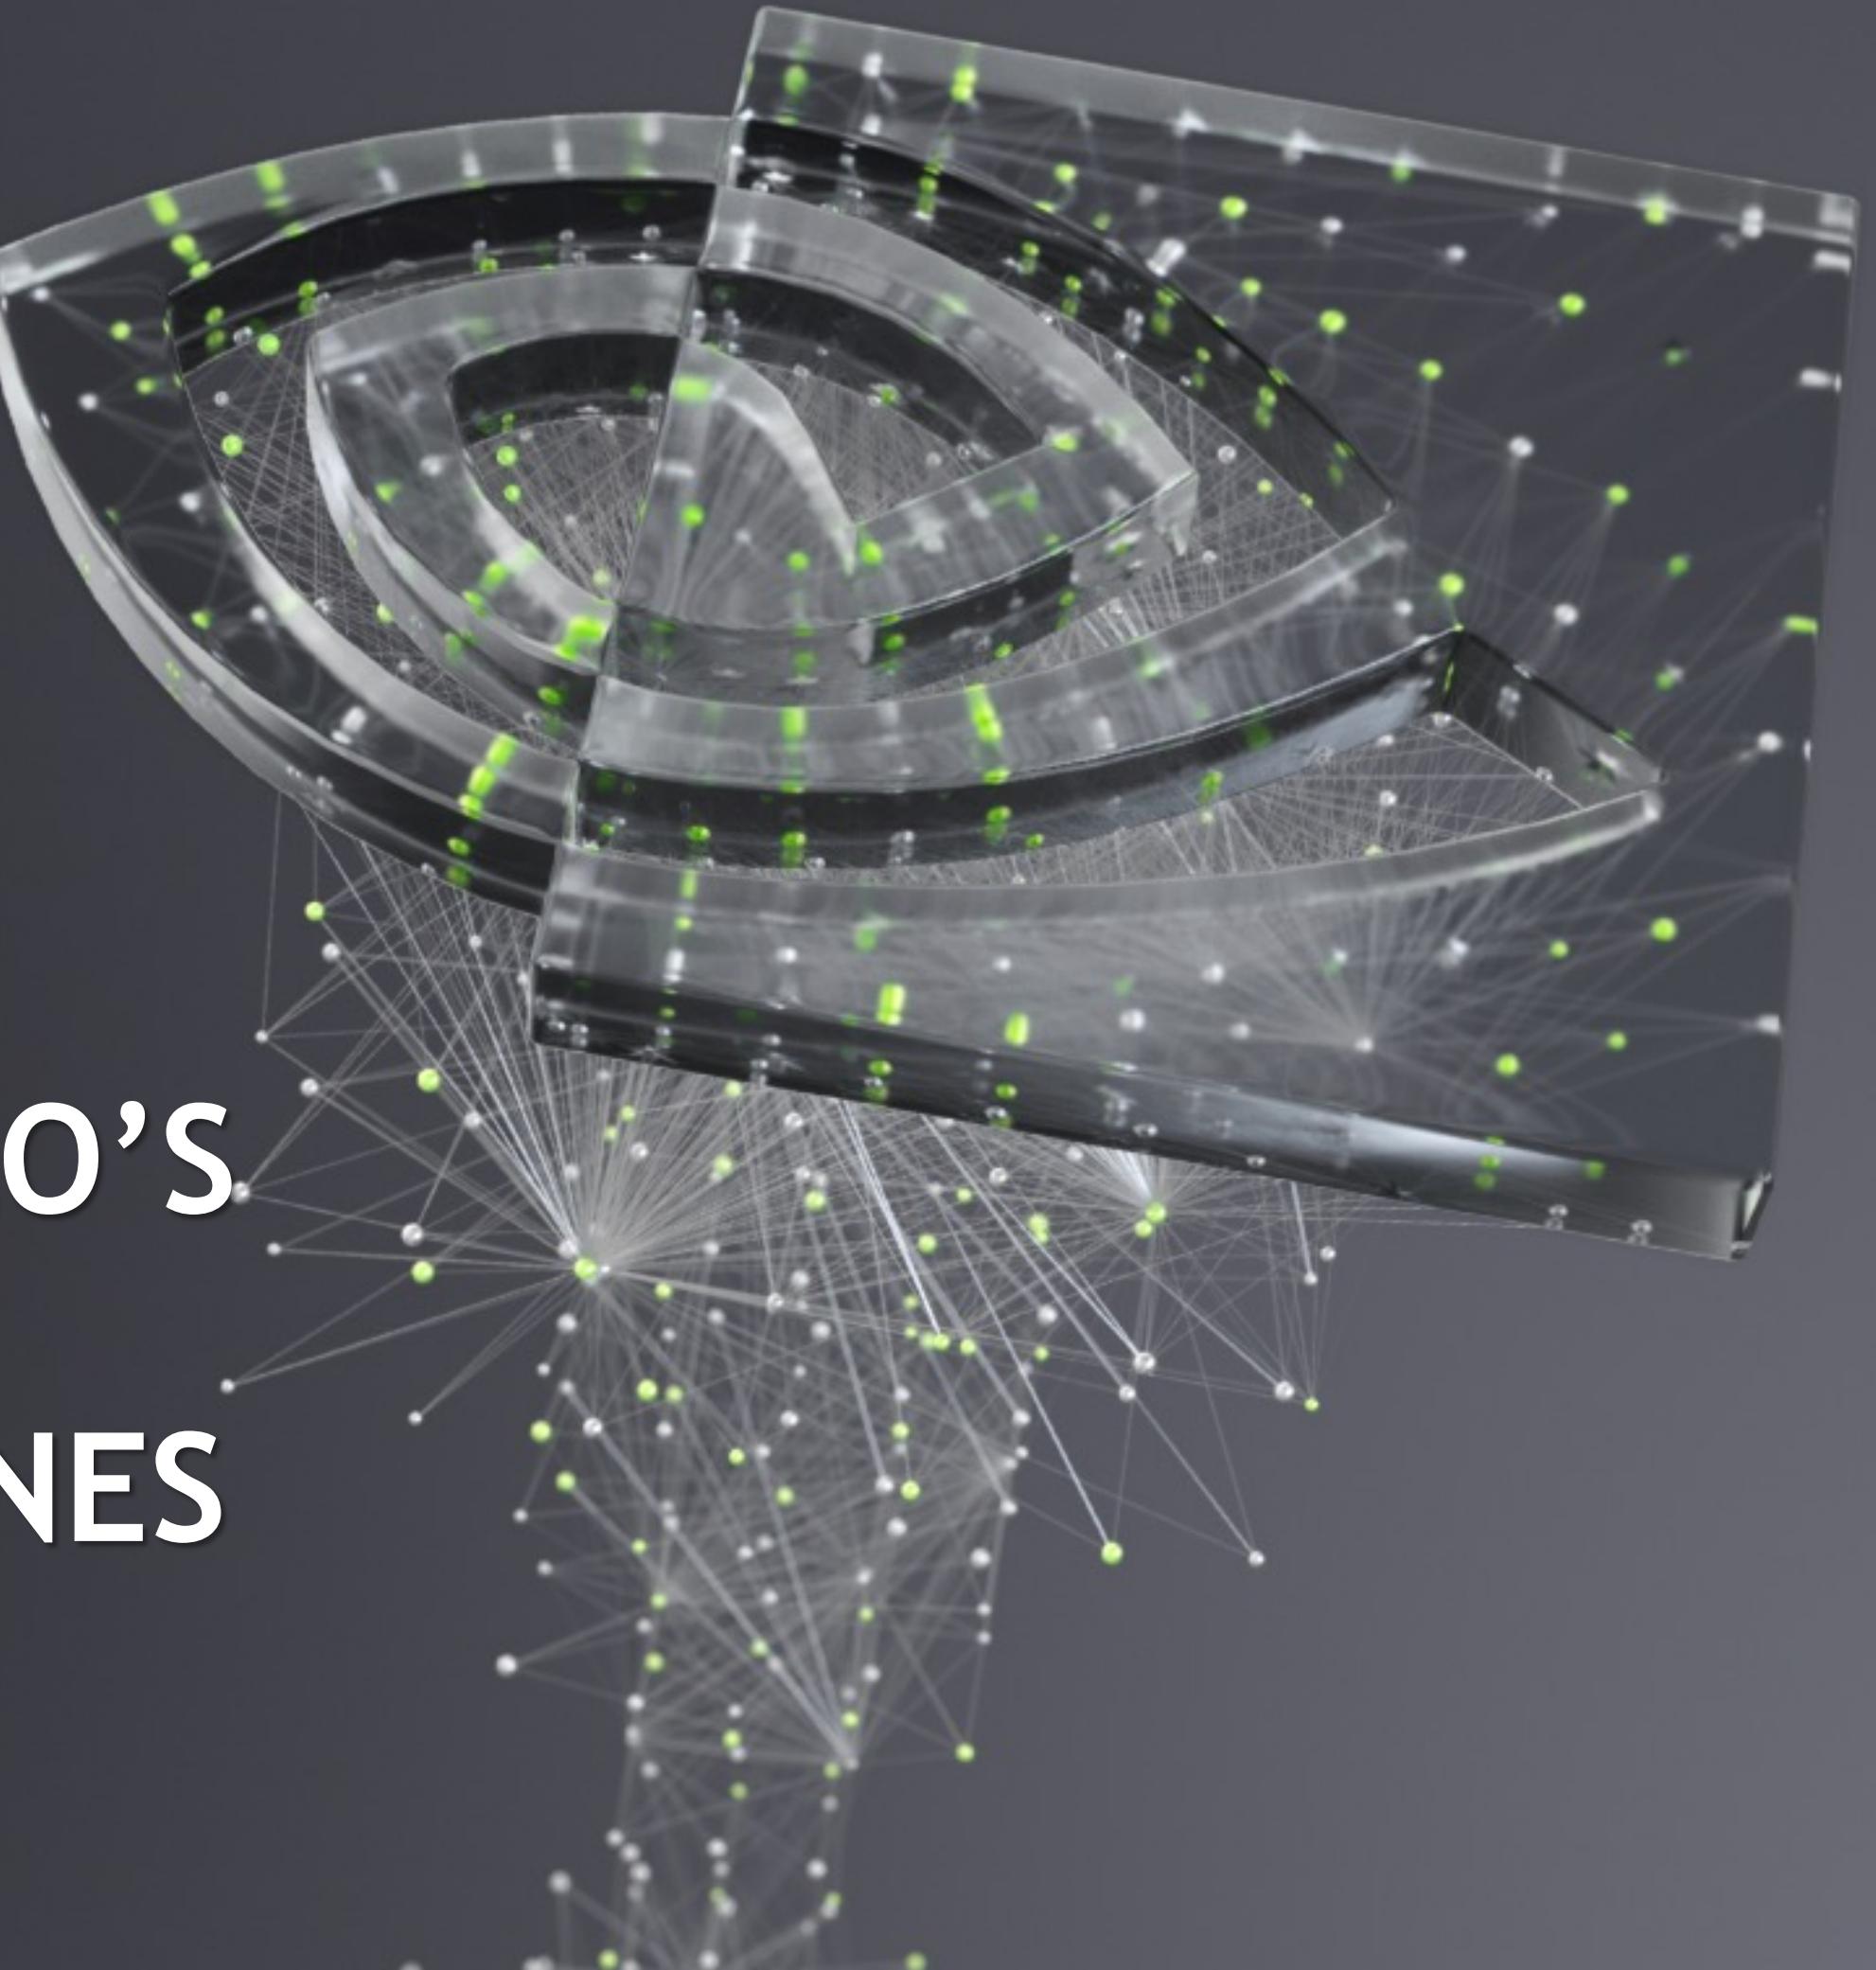




NVIDIA®

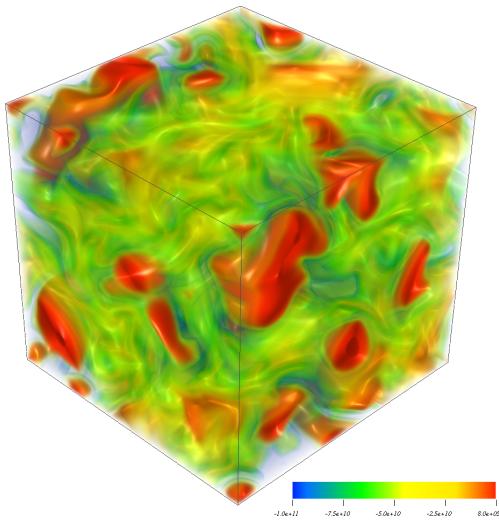
# REPRISING TOMASULO'S ALGORITHM FOR DISTRIBUTED MACHINES

Michael Bauer, 10/21/25



# HOW PROGRAMS RUN ON OUR DESKTOPS

Building applications is “easy”



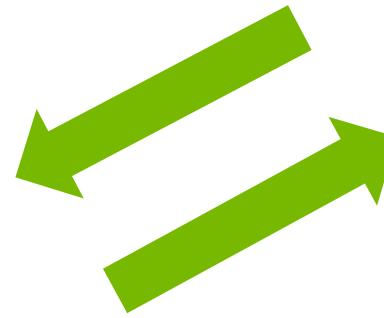
Direct Numerical Simulation of  
3D Combustion with Multi-Physics Chemistry



Cells: 4D Arrays: [X,Y,Z,Chemical Species]  
Fluxes: 5D Arrays [X,Y,Z,Chemical Species,Face]

```
cells[:, :] += alpha * np.sum(fluxes, axis=4)
```

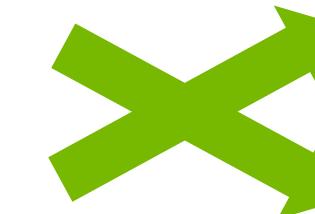
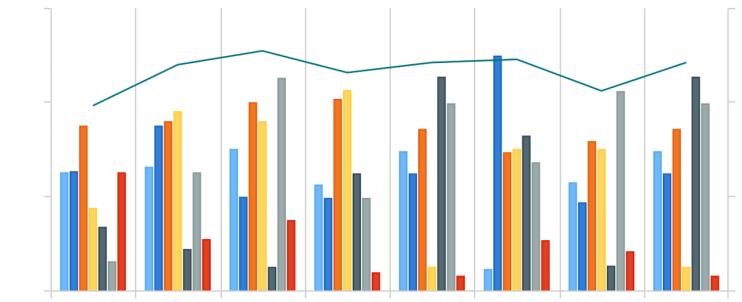
Uncertainty  
Quantification



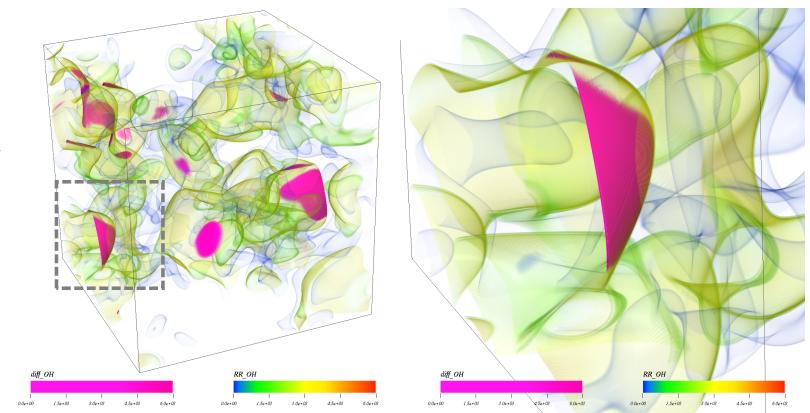
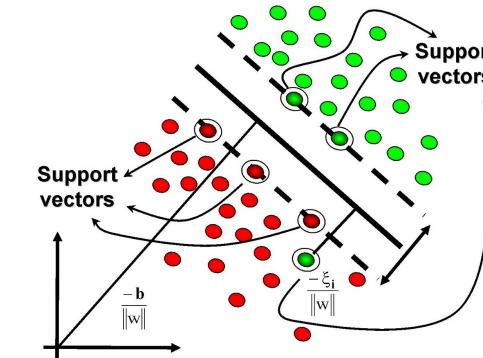
Statistical Analyses  
Group-by over chemical  
species with different  
compositions



**matplotlib**



SVM Classifier  
Recognize flame wavefronts

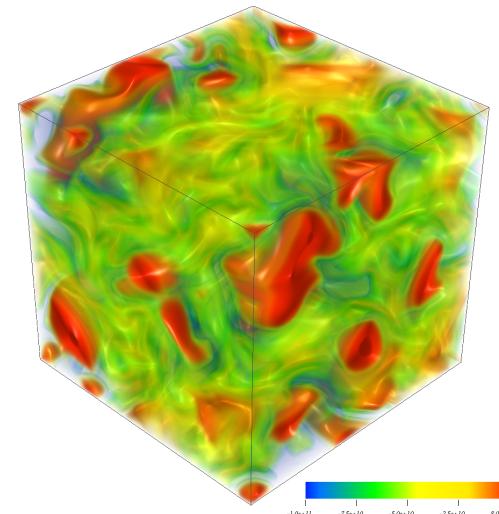


# HOW PROGRAMS RUN ON OUR SUPERCOMPUTERS



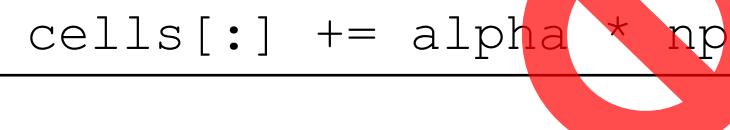
Enter mpi4py...  
Run on N nodes

Uncertainty  
Quantification



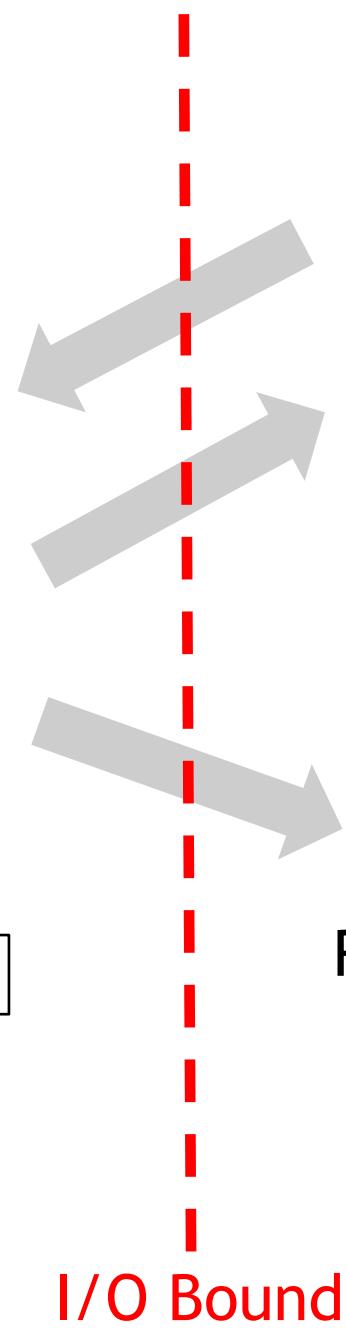
Manage explicit  
ghost cells

```
cells[:, :] += alpha * np.sum(fluxes, axis=4)
```



Building applications is “hard”

Write/Read  
to/from Disk



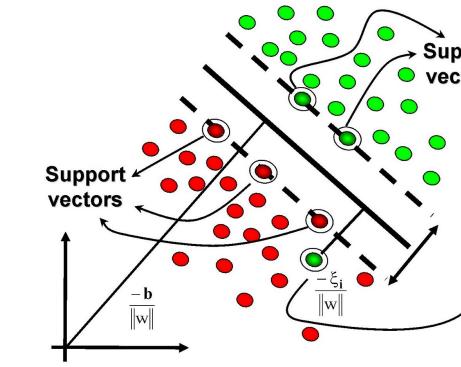
Run on M nodes



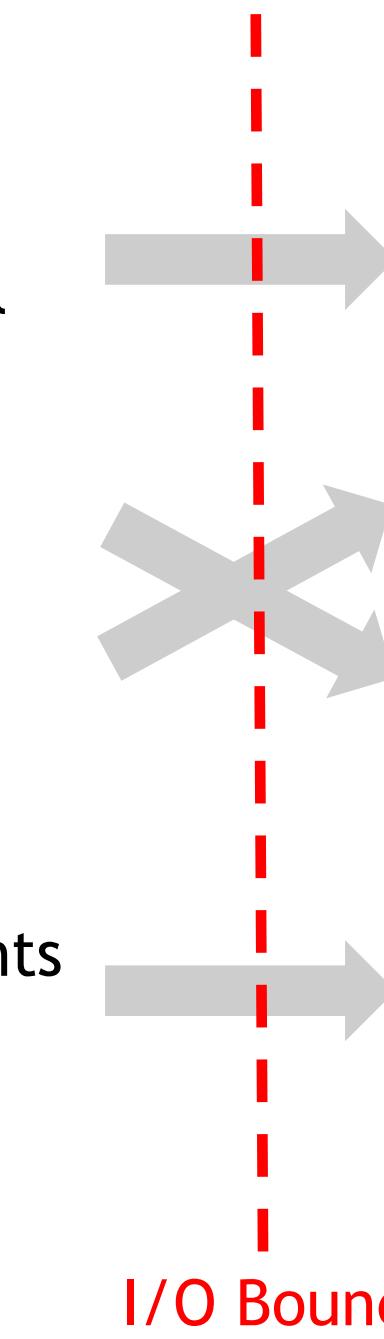
Statistical Analyses  
Group-by over chemical  
species with different  
compositions



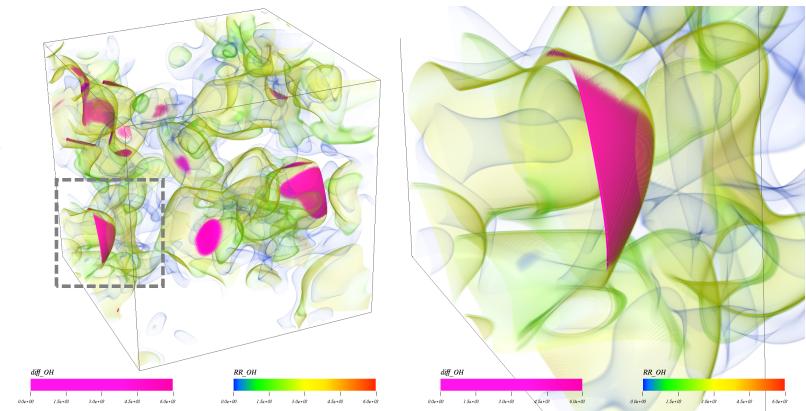
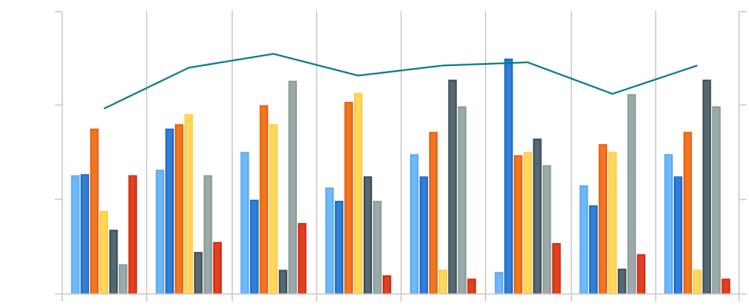
SVM Classifier  
Recognize flame wavefronts



Write/Read  
to/from Disk



Run on 1 node



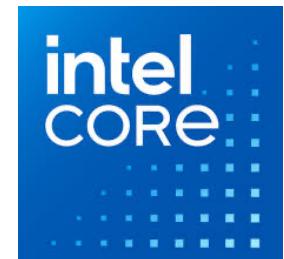
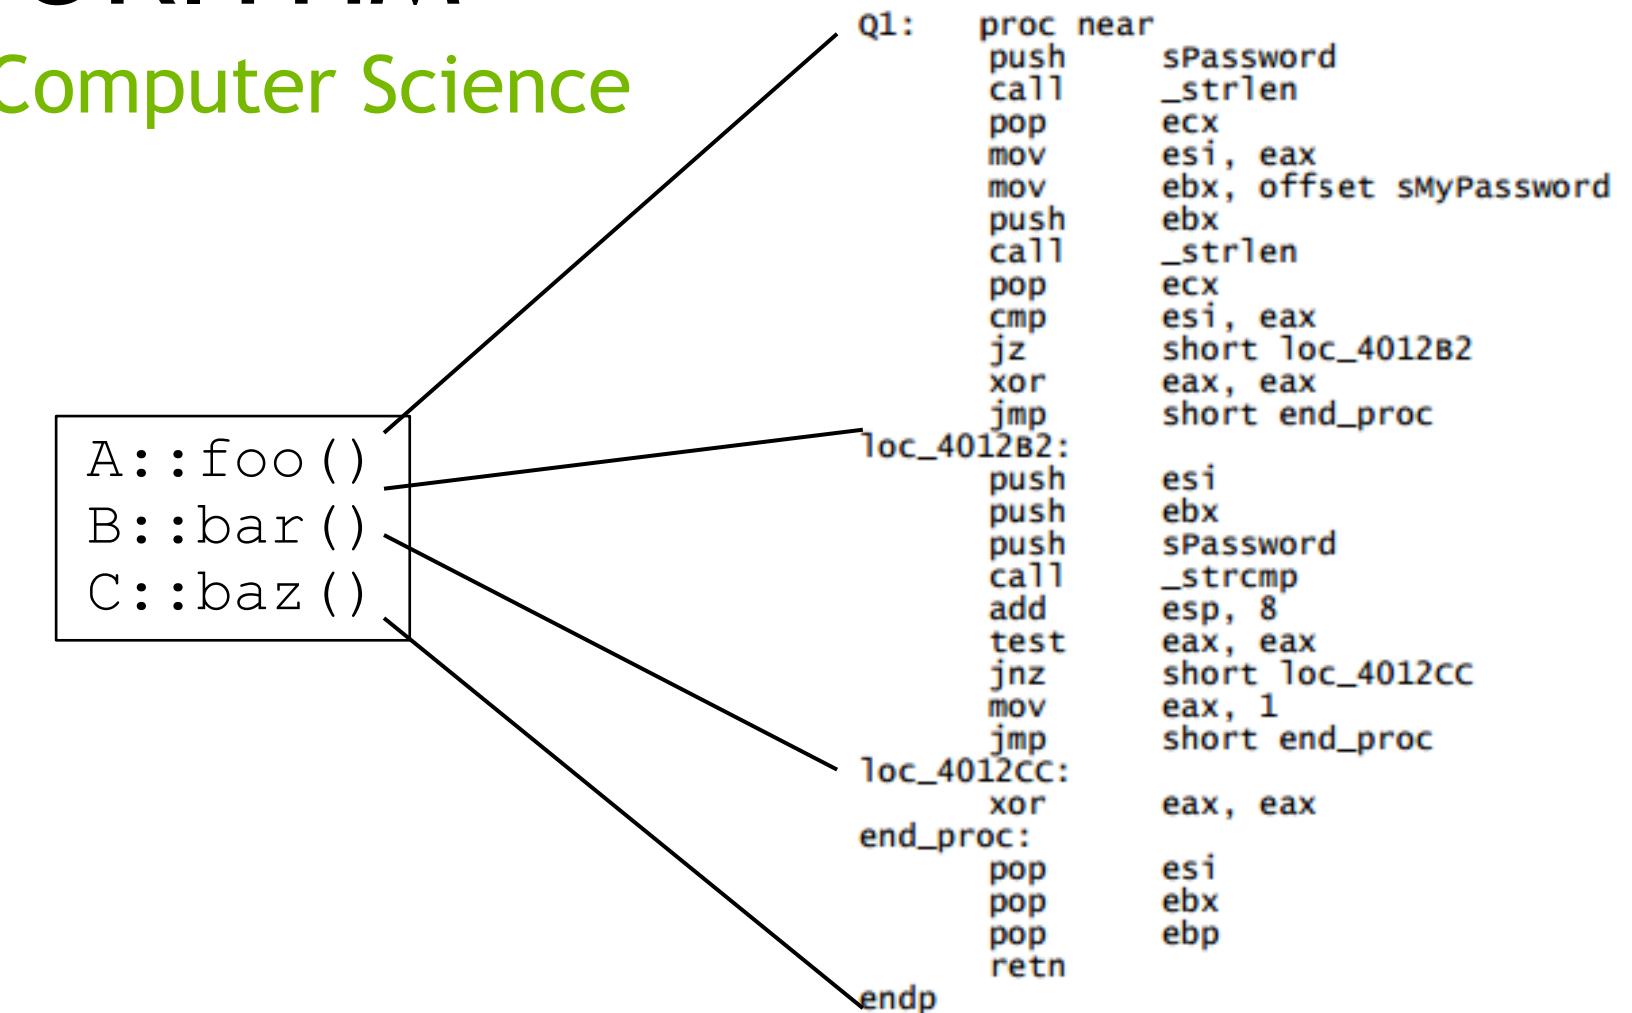
# TOMASULO'S ALGORITHM

One of the Best Abstractions in Computer Science

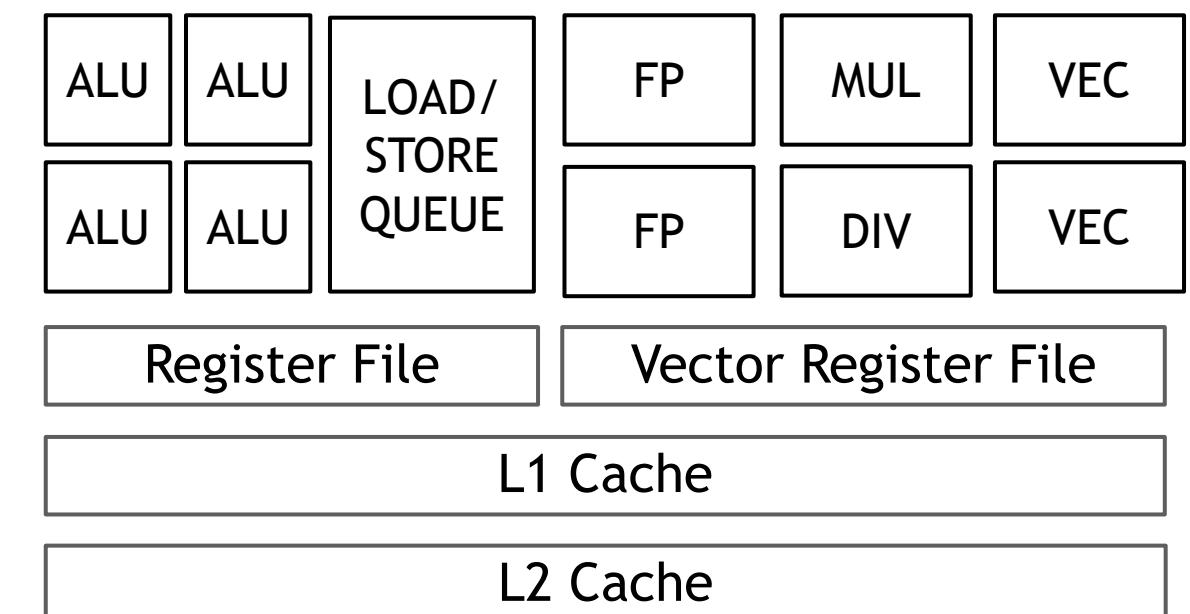
Hardware has lots of parallel functional units and yet neither programmers nor compilers have to think about these details

Benefits:

- Sequential semantics / implicit parallelism
- Automatic latency hiding
- Coherence of data
- Composition of software
- Precise exceptions
- Portability



Sequential Instruction Set Architecture (ISA)



# HARDWARE IS SELF-SIMILAR OVER SCALES

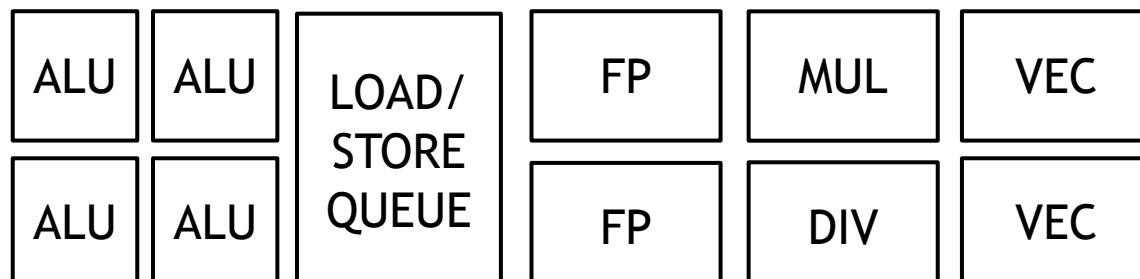
Why isn't software?

Increasing scale of hardware/computation

Processor



Sequential Instruction Set Architecture (ISA)



Register File      Vector Register File

L1 Cache

L2 Cache

Elegant Easy-to-Use Programming Model

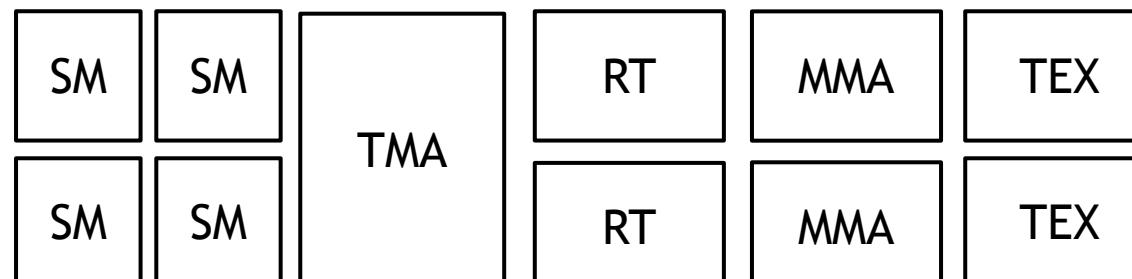
Sequential Semantics/Implicit Parallelism

Composable Software

Chip



???



SHARED MEMORY      TENSOR MEMORY

NOC

L2 Cache

Challenging Programming Models (MPI/CUDA/UPC/Co-Array Fortran/Global Arrays ...)

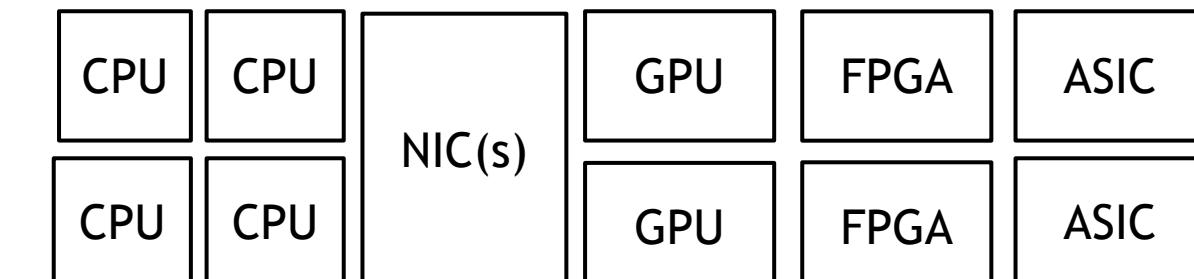
Explicit Parallelism, Explicit Data Movement, Explicit Synchronization

Hardly Any Software Works Together Without Massive Effort

Supercomputer



???



DRAM      HBMEM

NVLink/Infiniband/Slingshot

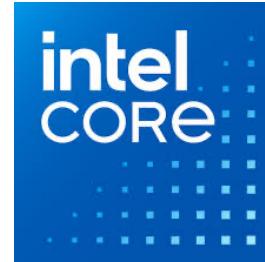
Distributed File System with Burst Buffers

# TOMASULO'S ALGORITHM STILL HAS MUCH TO OFFER

Remember the Lessons We Should Already Have Learned

Increasing scale of hardware/computation

Processor



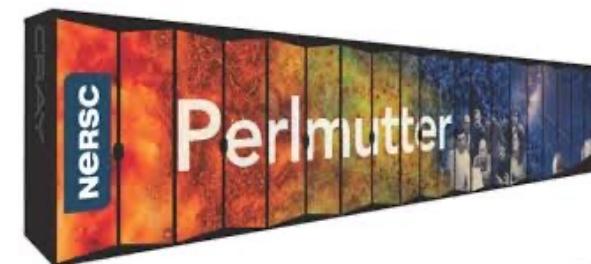
Sequential Instruction Set Architecture (ISA)

Tomasulo's Algorithm in Hardware

Chip



Supercomputer



Implicitly Parallel Task-Based Programming Models

Runtime Systems that Implement Tomasulo's Algorithm

Programs Are Dynamic Sequence of Tasks

Hygienic and ergonomic software ecosystems demand that (almost) all parallelism should be implicit

Tasks Name Region Usage

All software should trivially compose regardless of scale

Tasks Recursively Decompose At Different Scales

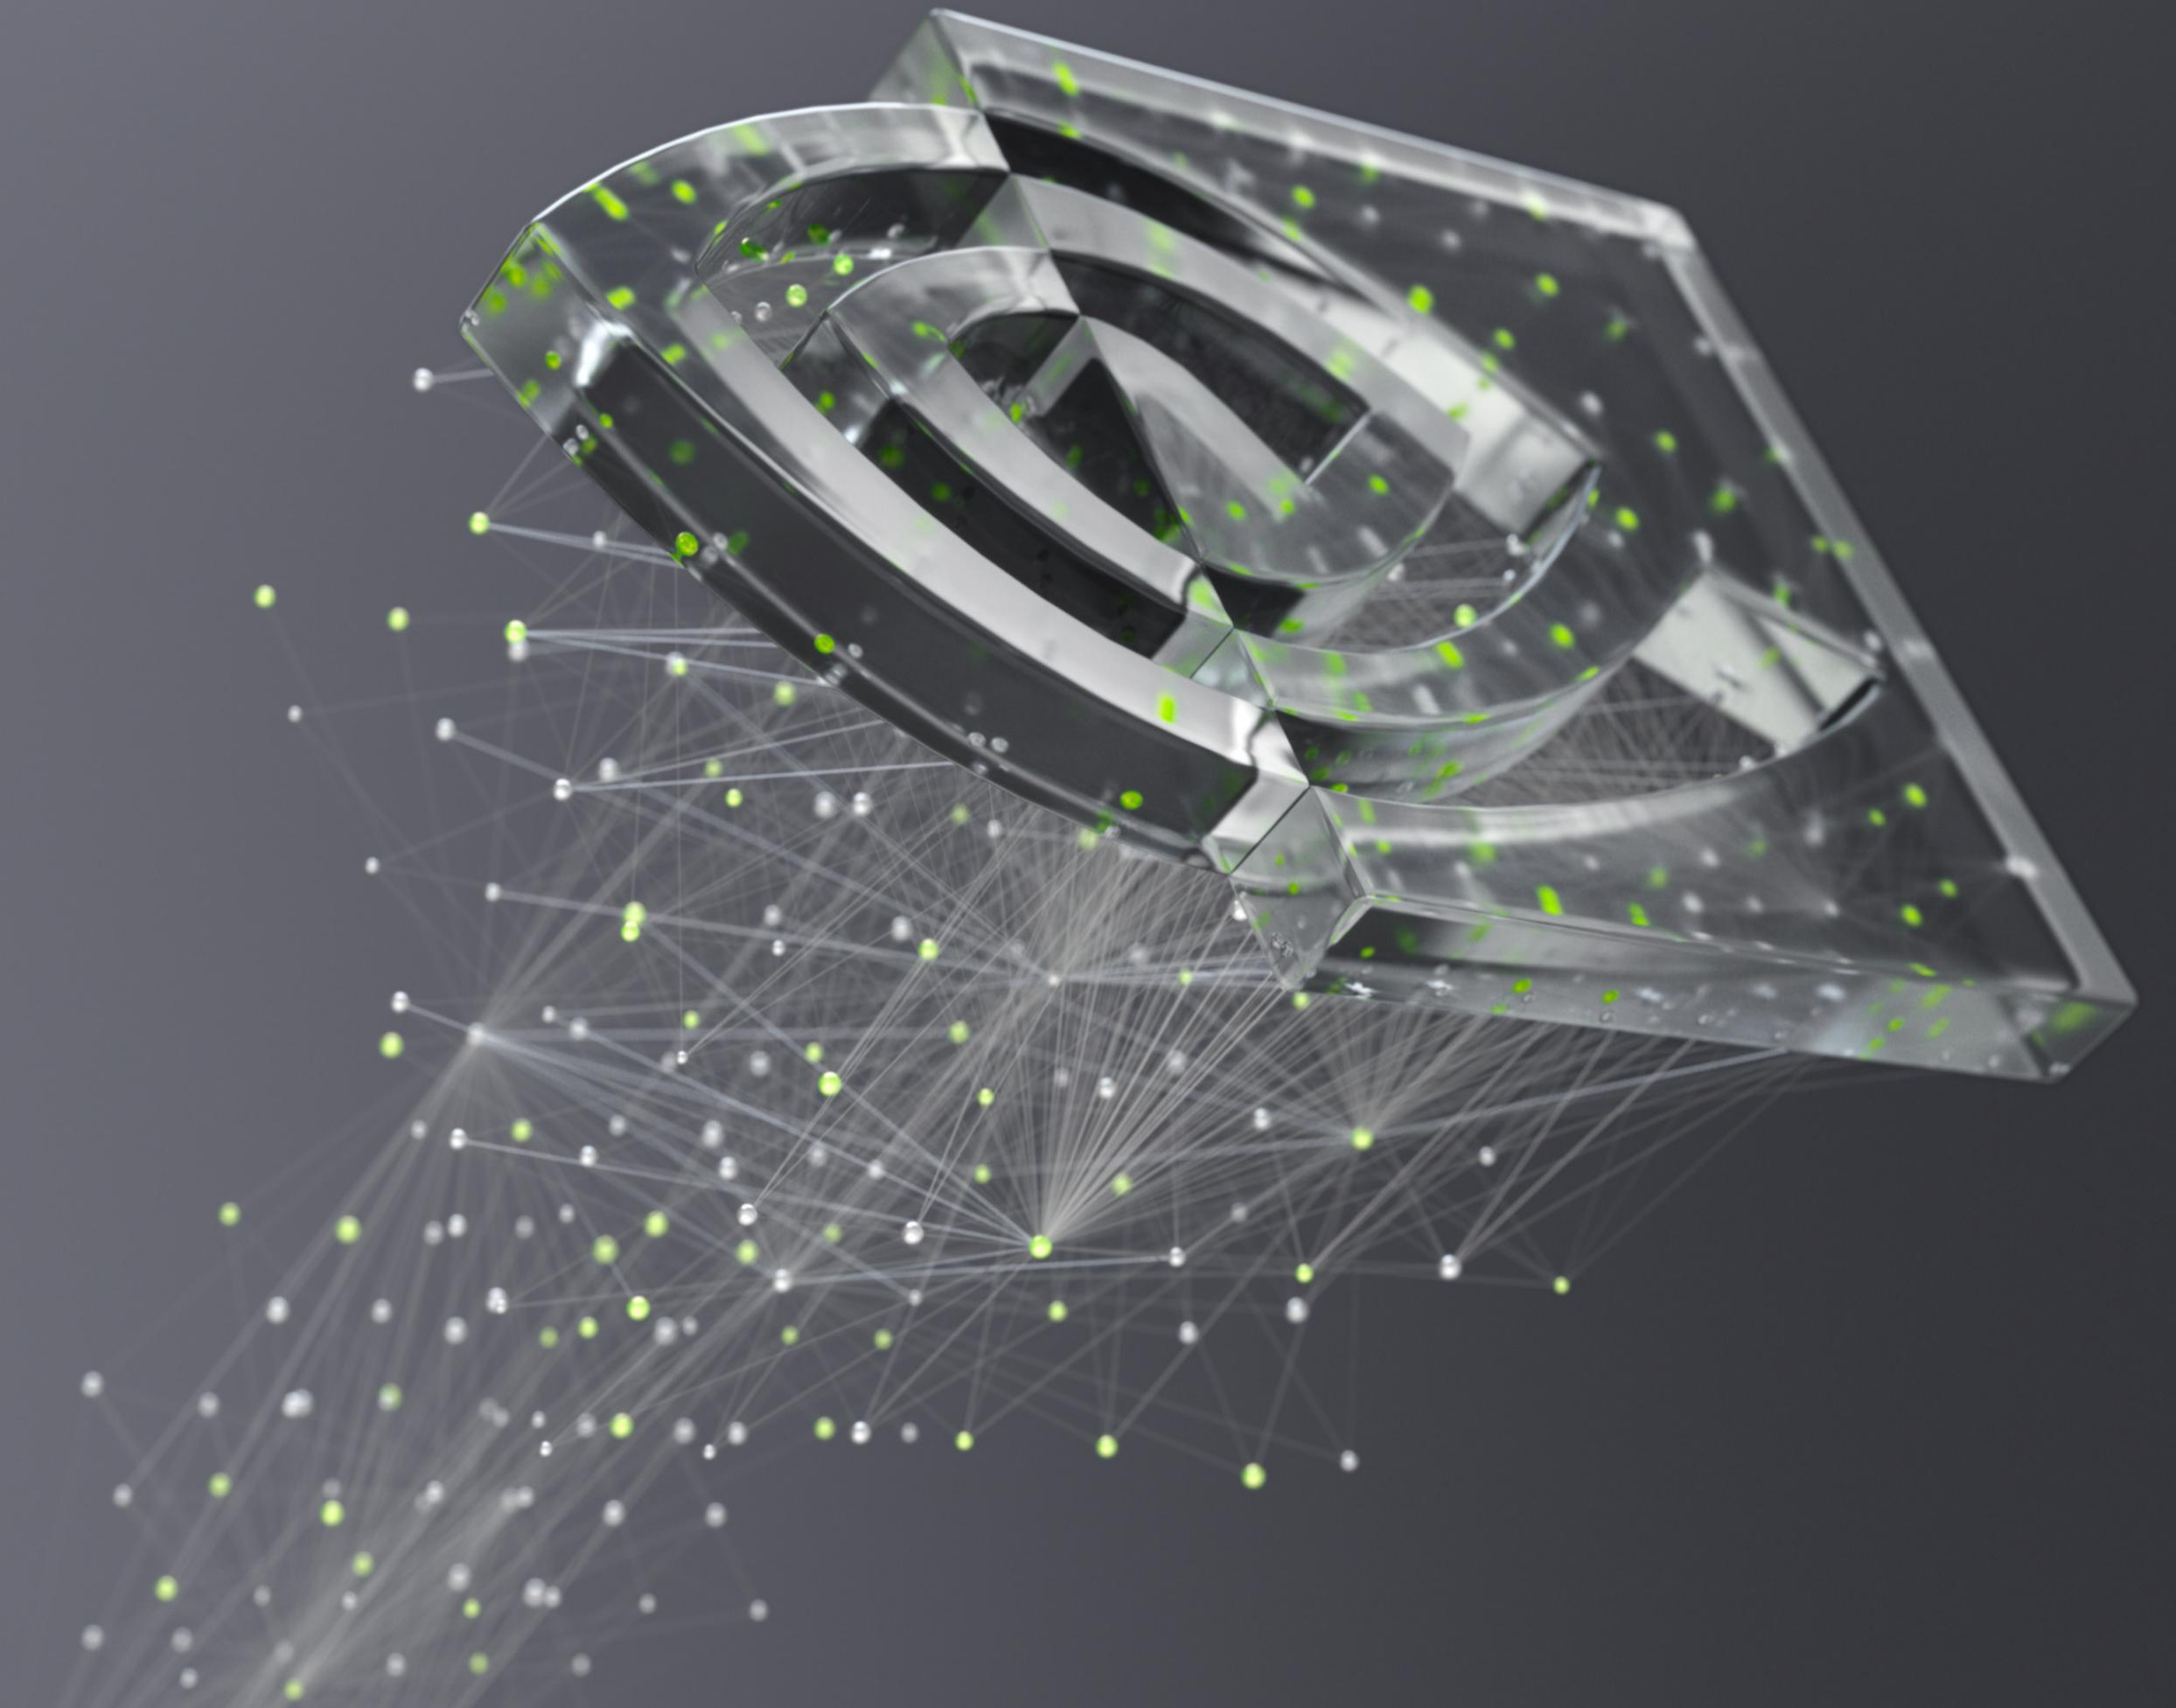
All software should be portable

t1: r1, r3, r4  
t2: r2, r4  
t3: r1, r2, r5, r6  
t4: r3, r4, r6  
t5, r2, r4, r5

Lack of discipline on our part to maintain good abstractions like those enabled by Tomasulo's algorithm

ta: r1, r5  
tb: r2, r5, r6  
tc: r2, r6  
td: r1, r2

Instructions : Registers :: Tasks : Regions



NVIDIA®