



Jeff Hammond  [@science\\_dot](https://c.im/@jeffscience)

Replying to @science\_dot, @miguelraz\_ and @JuliaLanguage

Julia is of course great because it's basically Fortran for people who are too lazy to declare types and has an interpreter.

1:41 pm · 22 Feb 2023 · 1,257 Views

---

1 Quote Tweet    22 Likes

# Why I think Julia is THE Language for Scientific Computing

A language should not “force you to rethink the way you solve problems”

Not everything needs to be a f-ing for loop!

Don't keep secrets from your compiler

OpenMP version:

```
// Begin the target region for GPU execution
#pragma omp target teams distribute parallel for \
    reduction(+:d_sum) \
    map(to: d_arr[0:N])
for (int i = 0; i < N; i++) {
    d_sum += d_arr[i] * d_arr[i];
}
```

Julia version:

```
d_squared = map(x -> x^2, d_arr)
d_sum = reduce(+, d_squared)
```

Pipelined Julia version:

```
d_sum = map(x -> x^2, d_arr) |> reduce(+)
```



# Why I think Julia is THE Language for Scientific Computing

A language should not “force you to rethink the way you solve problems”

There is no “Julian” way of programming

Not everything needs to be a f-ing for loop!

Julia natively contains structures, and concepts for parallel and pipelined work

Don’t keep secrets from your compiler

Multiple dispatch, JIT, LLVM, and introspection tools make Julia a powerful “LLVM frontend”

3

OpenMP version:

```
// Begin the target region for GPU execution
#pragma omp target teams distribute parallel for \
    reduction(+:d_sum) \
    map(to: d_arr[0:N])
for (int i = 0; i < N; i++) {
    d_sum += d_arr[i] * d_arr[i];
}
```

Julia version:

```
d_squared = map(x -> x^2, d_arr)
d_sum = reduce(+, d_squared)
```

Pipelined Julia version:

```
d_sum = map(x -> x^2, d_arr) |> reduce(+)
```



# A Template System that Allows you to Focus on Science

using LinearAlgebra

```
loss(w,b,x,y) = sum(abs2, y - (w*x .+ b)) / size(y,2)
loss∇w(w, b, x, y) = ... These don't have to be array-ish functions (e.g. Numpy), they
lossdb(w, b, x, y) = ... can contain if, for, etc statements.
```

```
function train(w, b, x, y ; lr=.1)
    w -= lmul!(lr, loss∇w(w, b, x, y))
    b -= lr * lossdb(w, b, x, y)
    return w, b
end

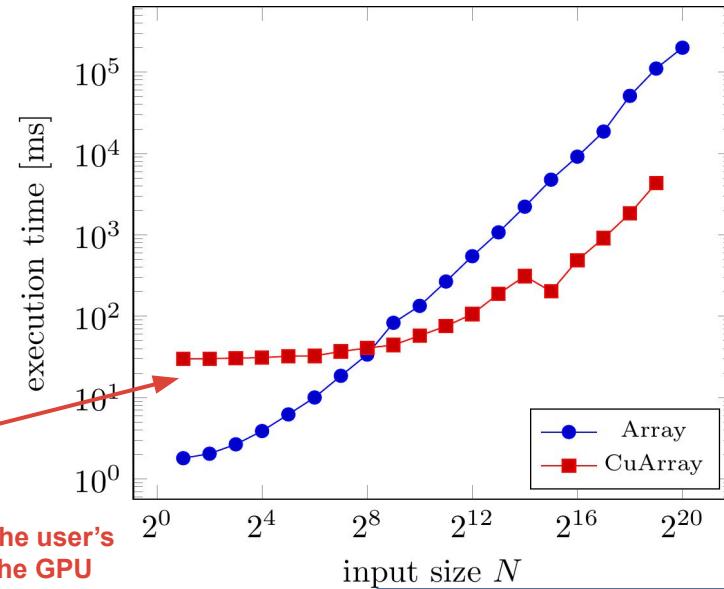
n = 100; p = 10
x = randn(n,p)'
y = sum(x[1:5, :]; dims=1) .+ randn(n)'*0.1
w = 0.0001*randn(1,p)
b = 0.0

for i=1:50
    w, b = train(w, b, x, y)
end
```

x = CuArray(x)  
y = CuArray(y)  
w = CuArray(w)

By moving your data to the GPU, the user's of the algorithm "Just works" on the GPU

Rapid software prototyping for heterogeneous and distributed platforms  
Besard T., Churavy V., Edelman A., De Sutter B.  
(doi:10.1016/j.advengsoft.2019.02.002)



# A Template System that Allows you to Focus on Science

Abstraction, Specialization, and Multiple Dispatch

## 1. Abstraction to obtain generic behavior:

Encode behavior in the type domain:

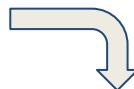
```
transpose(A::Matrix{Float64})::Transpose{Float64, Matrix{Float64}}
```

Did I really need to move  
memory for that transpose?

## 2. Specialization of functions to produce optimal code

## 3. Multiple-dispatch to select optimized behavior

```
rand(N, M) * rand(K, M)'  
Matrix * Transpose{Matrix}  
  
function mul!(C::Matrix{T}, A::Matrix{T}, tB::Transpose{<:Matrix{T}}, a, b) where {T<:BlasFloat}  
    gemm_wrapper!(C, 'N', 'T', A, B, MulAddMul(a, b))  
end
```



compiles to

No I did not! I know  $AB^T$  is  
the dot product of every row  
of A with every row of B .

```

using CUDA
function gemm!(A,B,C)
    row = (blockIdx().x - 1) * blockDim().x + threadIdx().x
    col = (blockIdx().y - 1) * blockDim().y + threadIdx().y
    sum = zero(eltype(C))

using Metal
function gemm!(A,B,C)
    row, col = thread_position_in_grid_2d()

    sum = zero(eltype(C))

    if row <= size(A, 1) && col <= size(B, 2)
using AMDGPU
function gemm!(A,B,C)
    row = (workgroupIdx().x - 1) * workgroupDim().x + workitemIdx().x
    col = (workgroupIdx().y - 1) * workgroupDim().y + workitemIdx().y
    sum = zero(eltype(C))

    if row <= size(A, 1) && col <= size(B, 2)
        for i = 1:size(A, 2)
            @inbounds sum += A[row, i] * B[i, col]
using oneAPI
function gemm!(A,B,C)
    row = get_global_id(0)
    col = get_global_id(1)

    sum = zero(eltype(C))

    if row <= size(A, 1) && col <= size(B, 2)
        for i = 1:size(A, 2)
            @inbounds sum += A[row, i] * B[i, col]
        end
        @inbounds C[row, col] = sum
    end

    return
end

```

- Leverage Julia's Toolchain for low-level portable code!
- Julia's xPU LLVM-backend also can generate Kernel code using CUDA.jl, AMDGPU.jl, or oneAPI.jl
- KernelAbstractions.jl wraps all of these up into a single @kernel macro



```

using KernelAbstractions
@kernel function gemm!(A, B, C)
    row, col = @index(Global, NTuple)

    sum = zero(eltype(C))
    for i = 1:size(A, 2)
        @inbounds sum += A[row, i] * B[i, col]
    end
    @inbounds C[row, col] = sum
end

```

# Bonus slides



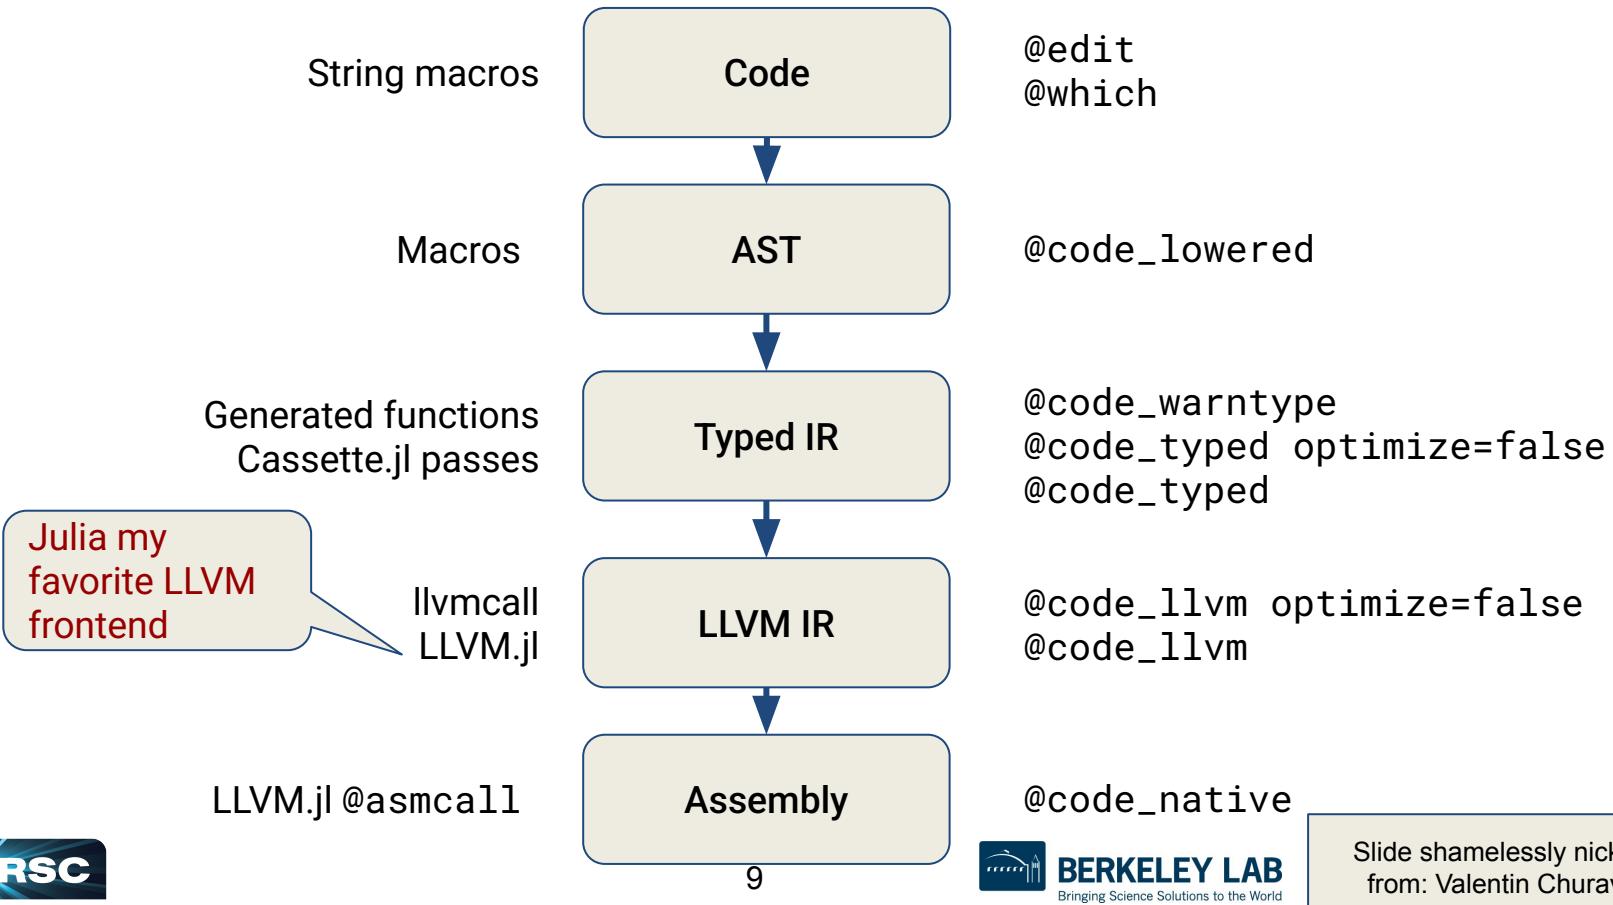
# Beware of the Glue-Code Pattern

- Some languages are not ABI-compatible
  - Calling from one language to another will incur a cost due to type conversion
- This will can get very expensive if calls come from the “inner loop” (eg. AI inference call for every grid cell)

Function signature	Pybind11	ccall		speedup	
int fn0()	132	$\pm 14.9$	2.34	$\pm 1.24$	56×
int fn1(int)	217	$\pm 20.9$	2.35	$\pm 1.33$	92×
double fn2(int, double)	232	$\pm 11.7$	2.32	$\pm 0.189$	100×
char* fn3(int, double, char*)	267	$\pm 28.9$	6.27	$\pm 0.396$	42×

all times in ns

# Introspection and staged metaprogramming



# Julia is basically a REPL for LLVM

Julia provides interfaces to the LLVM backend.

Eg.:

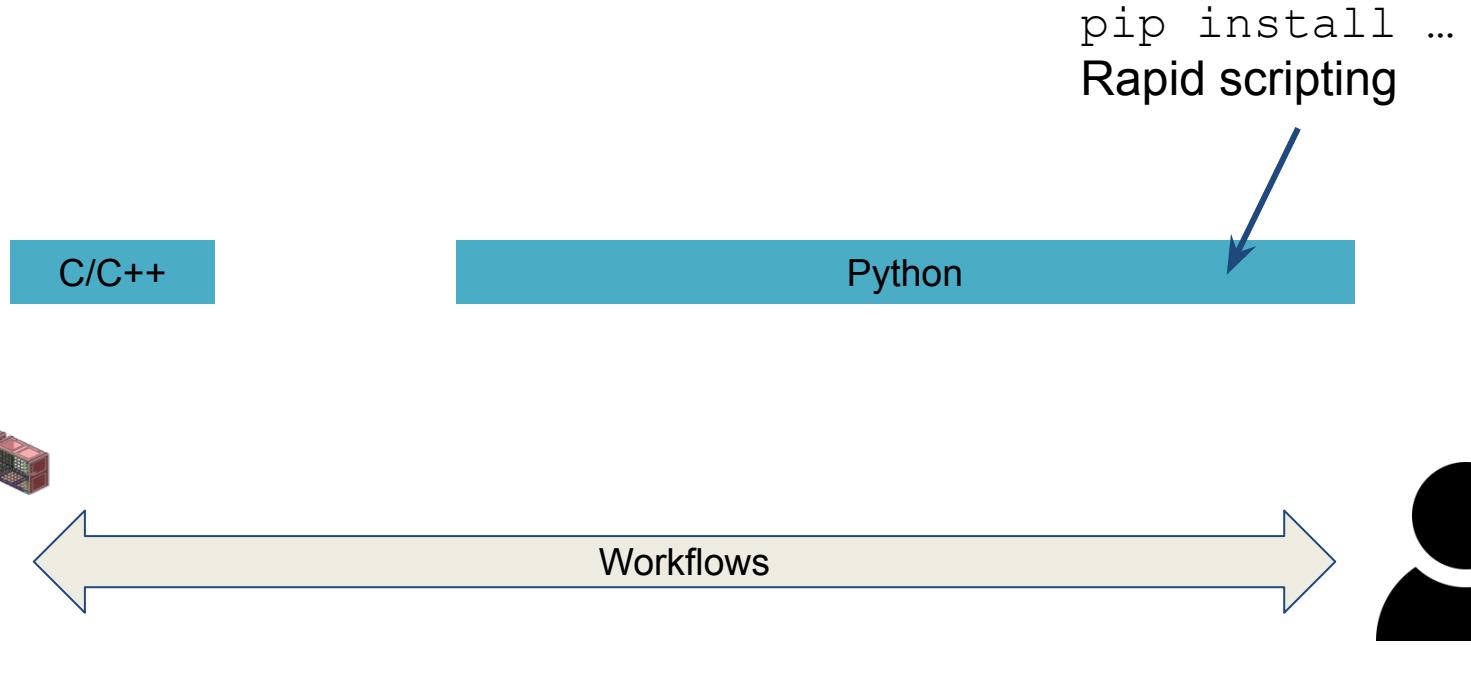
- loopinfo
- llvmpcall

```
[16]: macro unroll(expr)
    expr = loopinfo("@unroll", expr, (Symbol("llvm.loop.unroll.full"),))
    return esc(expr)
end

for (jlf, f) in zip((:+, :*, :-), (:add, :mul, :sub))
    for (T, llvmT) in ((:Float32, "float"), (:Float64, "double"))
        ir = """
            %x = f$f contract nsz $llvmT %0, %1
            ret $llvmT %x
        """
        @eval begin
            # the @pure is necessary so that we can constant propagate.
            @inline Base.@pure function $jlf(a::$T, b::$T)
                Base.llvmpcall($ir, $T, Tuple{$T, $T}, a, b)
            end
        end
    end
    @eval function $jlf(args...)
        Base.$jlf(args...)
    end
end
```



# No Annoying “Paradigm Shifts”



# No “Nibbling Around The Edges”

Function signature	Pybind11	ccall	speedup		
int fn0()	132	$\pm 14.9$	2.34	$\pm 1.24$	56×
int fn1(int)	217	$\pm 20.9$	2.35	$\pm 1.33$	92×
double fn2(int, double)	232	$\pm 11.7$	2.32	$\pm 0.189$	100×
char* fn3(int, double, char*)	267	$\pm 28.9$	6.27	$\pm 0.396$	42×

